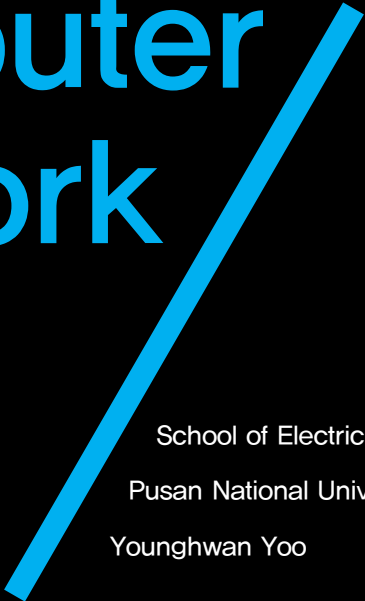


Computer Network

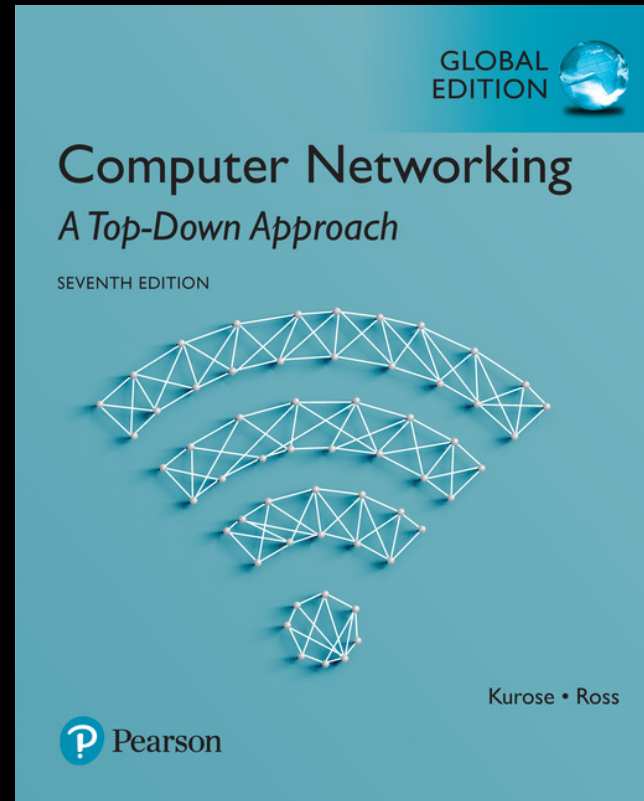


Transport Layer

School of Electric and Computer Engineering

Pusan National University, KOREA

Younghwan Yoo



Computer Networking

A Top-Down Approach

7th edition

Jim Kurose, Keith Ross

Pearson

April 2016

Contents

Computer Network introduction

01. Transport–Layer Services

02. Multiplexing and Socket

03. User Datagram Protocol

04. Reliable Data Transfer Principles

Contents

Computer Network introduction

05. Transmission Control Protocol

06. Congestion Control

07. TCP Congestion Control Algorithm

08. TCP vs. UDP



01. Transport-Layer Services

■ Program


- an executable file containing the set of instructions written to perform a specific job
- stored on a disk

■ Process

- an executing instance of a program
- resides on the primary memory
- several processes related to same program at the same time

■ Thread

- the smallest executable unit of a process




Transport layer

- logical communication between processes
- relies on, enhances, network layer services

Network layer

- logical communication between hosts



Send Side


- Application message **fragmentation** into segments
- Segment passing to network layer

Recv Side

- Segments **reassembly** into messages
- Message passing to application layer

Internet Transport Protocols

TCP, UDP



TCP

- Transmission Control Protocol
- **Reliable, in-order delivery**
- **Connection-oriented service**
 - connection setup
 - error control
 - flow control
 - congestion control

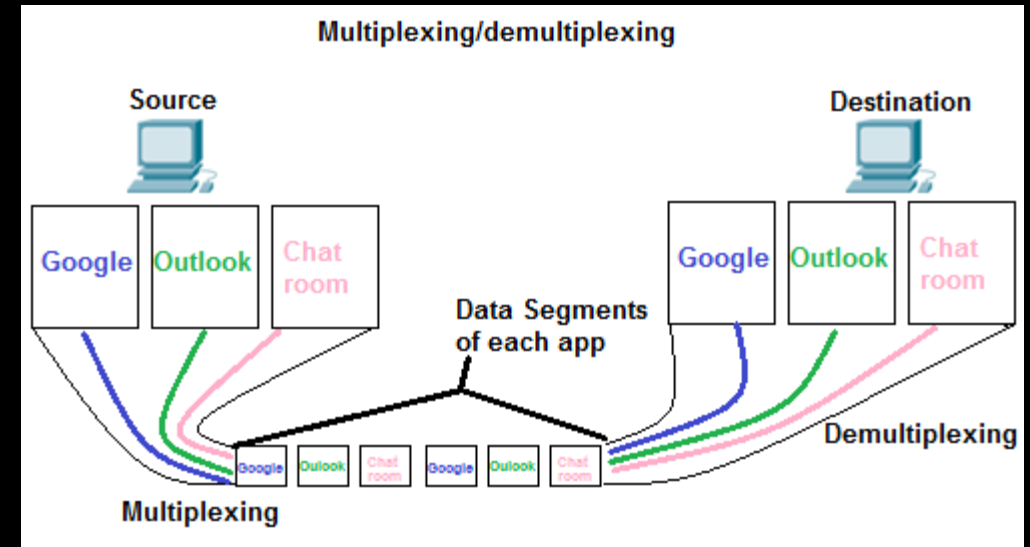
UDP

- User Datagram Protocol
- **Unreliable, unordered delivery**
- **Connectionless service**
 - faster than TCP



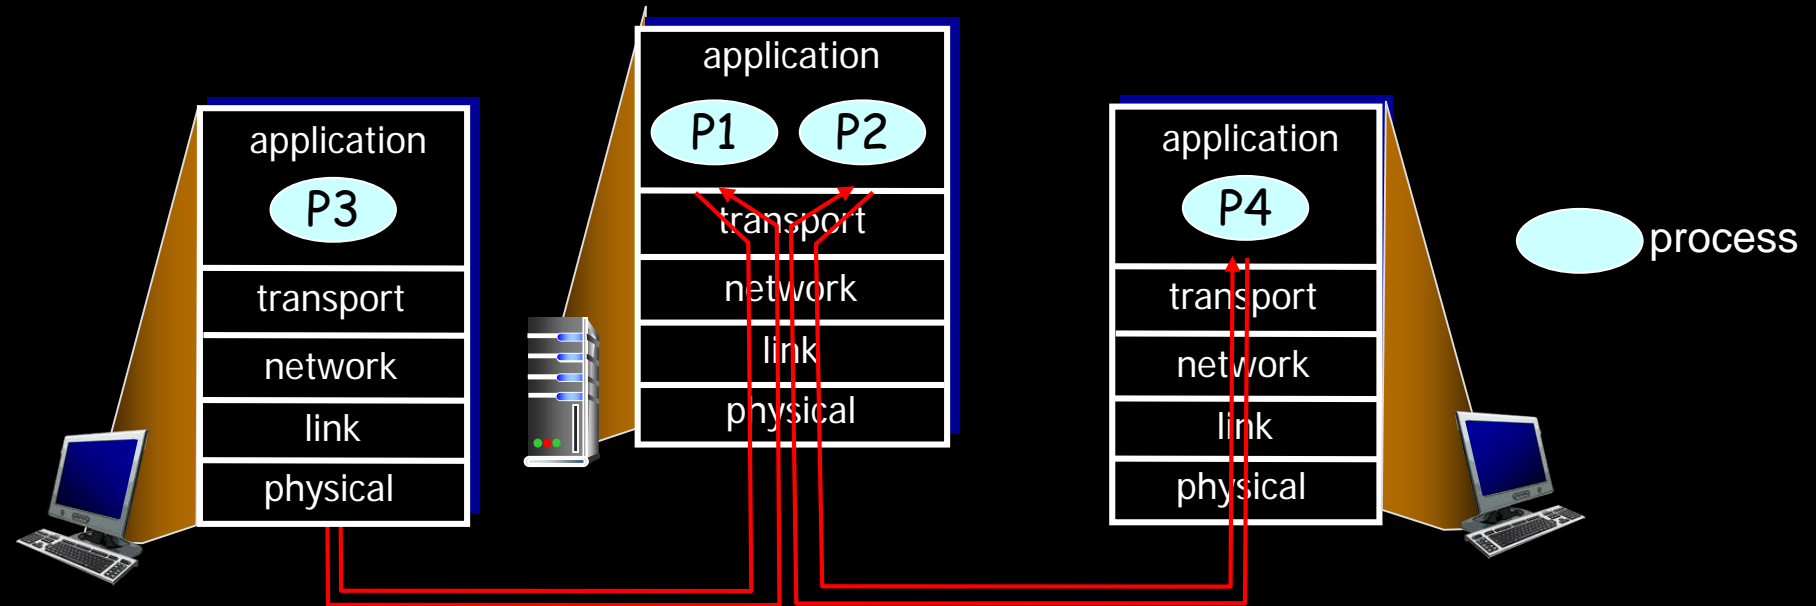
02. Multiplexing and Socket

- **Multiplexing** at sender
 - sending data from its own multiple applications through network
- **Demultiplexing** at receiver
 - delivering data packets to their appropriate receivers among its own multiple applications



출처 -

https://www.google.co.kr/url?sa=i&rct=j&q=&esrc=s&source=images&cd=&cad=rja&uact=8&ved=2ahUKewiNn_qt7ubbAhUU87wKHVnKCbQQjRx6BAgBEAU&url=http%3A%2F%2Fwww.cnt4all.com%2F2016%2F08%2F07-transport-layer-multiplexing-and.html&psig=AOvVaw03CLdYUqYYGWH6m155fosy&ust=1529740774014370

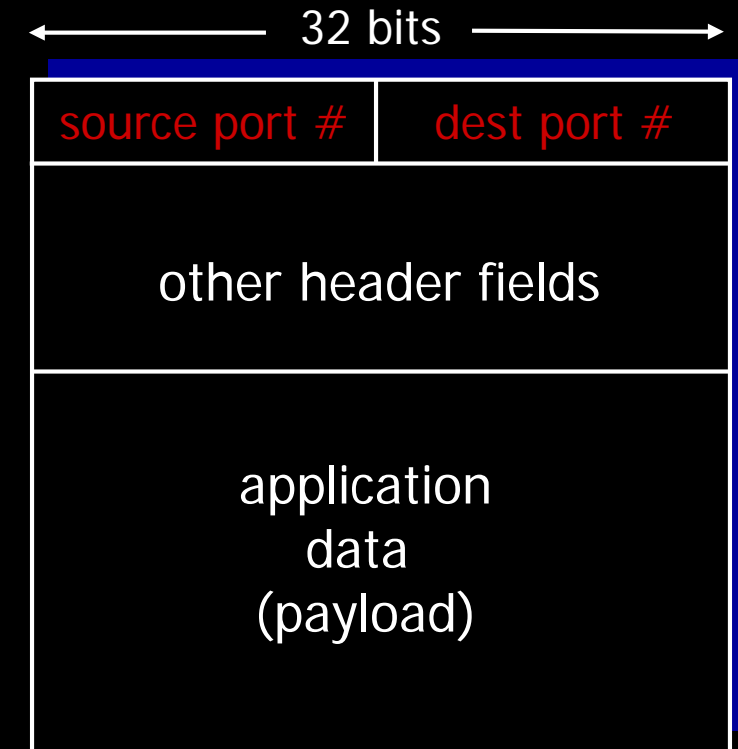


- The most basic role of the transport layer.

■ Port number

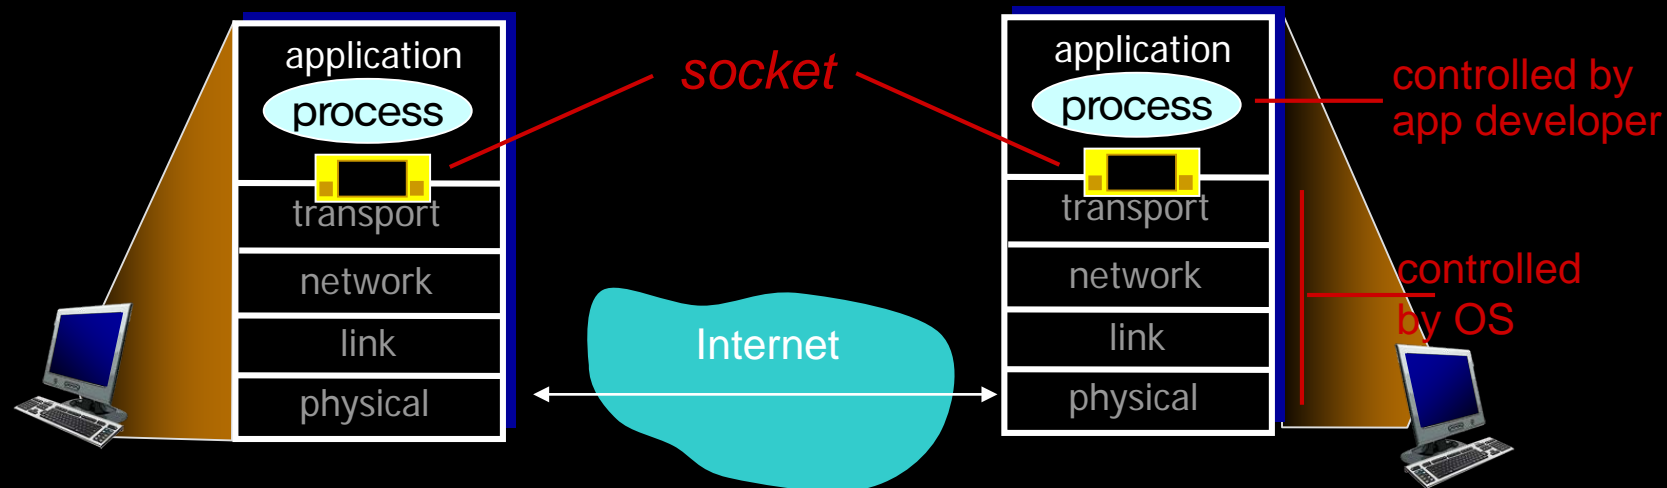
- different applications are assigned different port numbers
 - transport layer segments have fields for src/dst port numbers in common
 - used for differentiating segments
- However, the transport layer must be instructed by the application which process on which host should be the destination. How?

➡ *Socket !!!*



TCP/UDP segment format

- Process sends/receives messages to/from its **socket**
 - e.g., BSD socket, Winsock
- Analogous to door
 - sending process shoves message out door
 - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process



- The 'Socket' is provide as the form of APIs (application program interfaces) by OS

- Server

```
import socket

serversocket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
serversocket.bind(('127.0.0.1', 8089)) # localhost = 127.0.0.1
serversocket.listen(5)

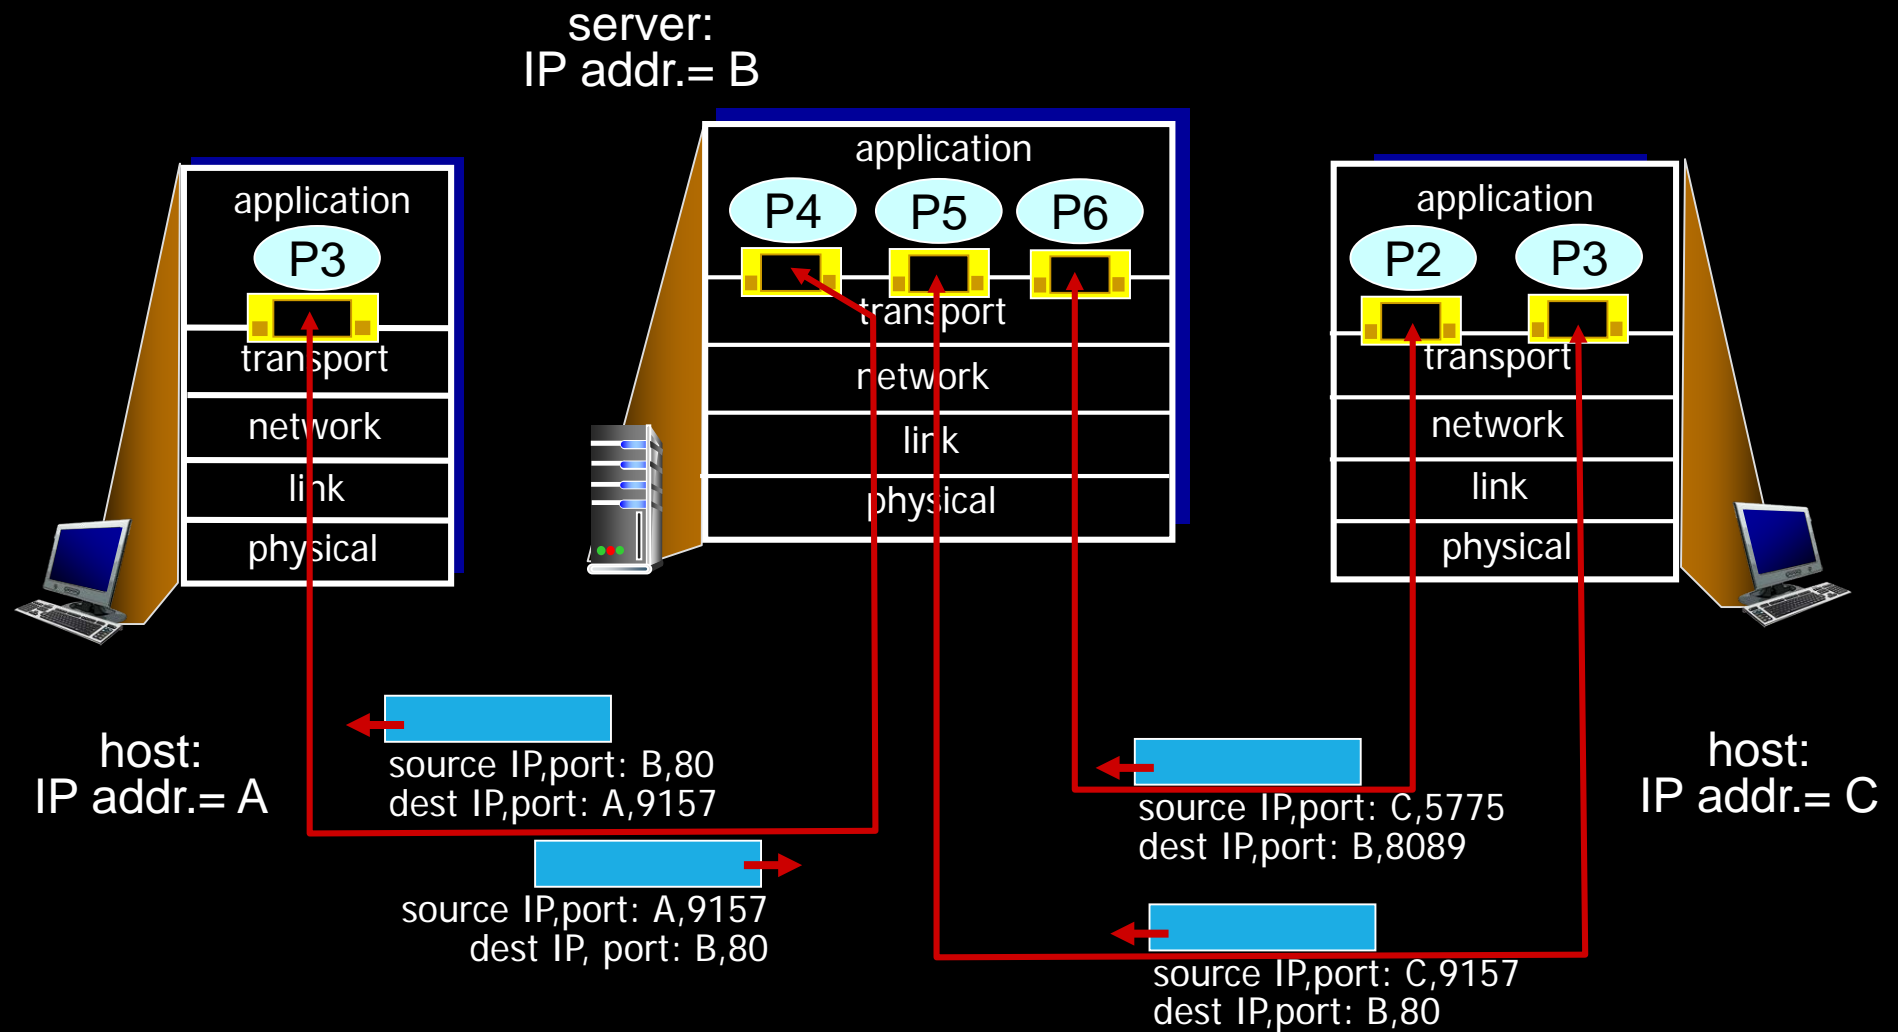
while True:
    connection, address = serversocket.accept()
    buf = connection.recv(64)
    if len(buf) > 0:
        print(buf)
        break
```

- Client

```
import socket

clientsocket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
clientsocket.connect(('localhost', 8089))
clientsocket.send(b'GET /index.html')
```

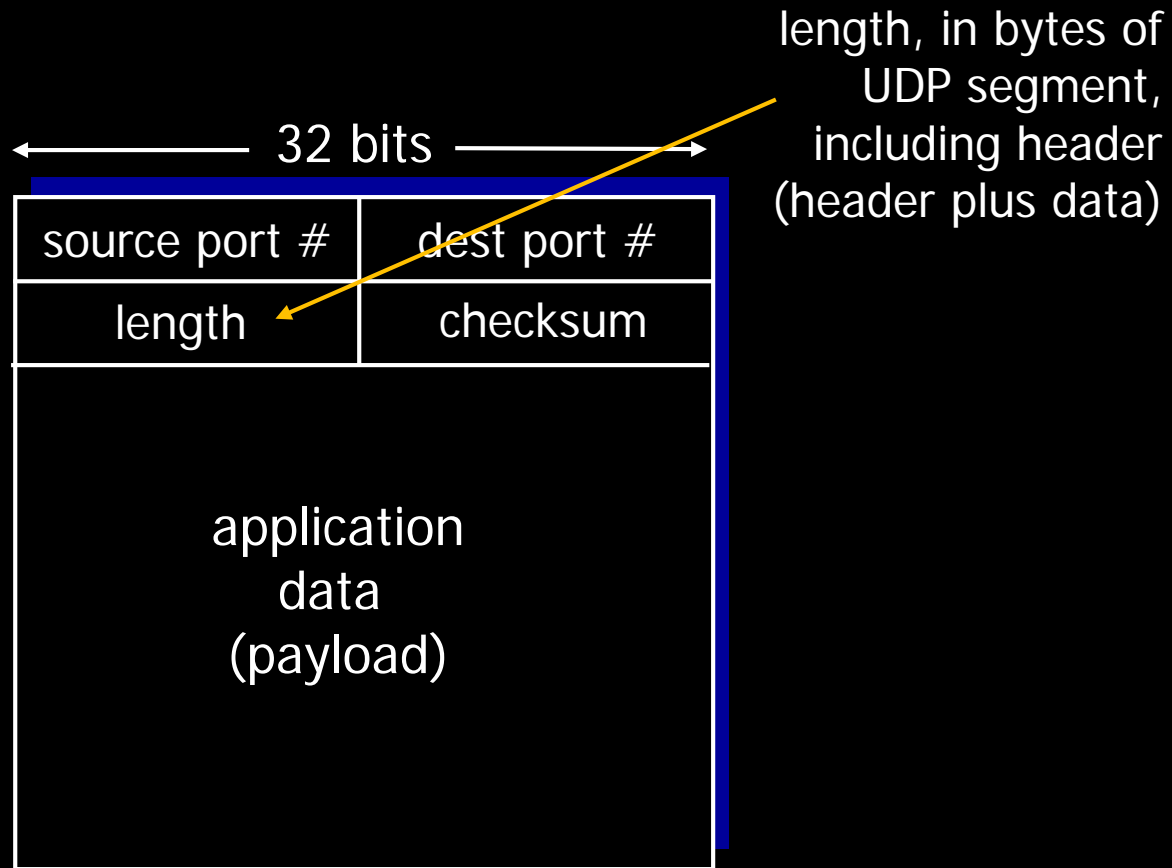
Demultiplexing Example





03. User Datagram Protocol

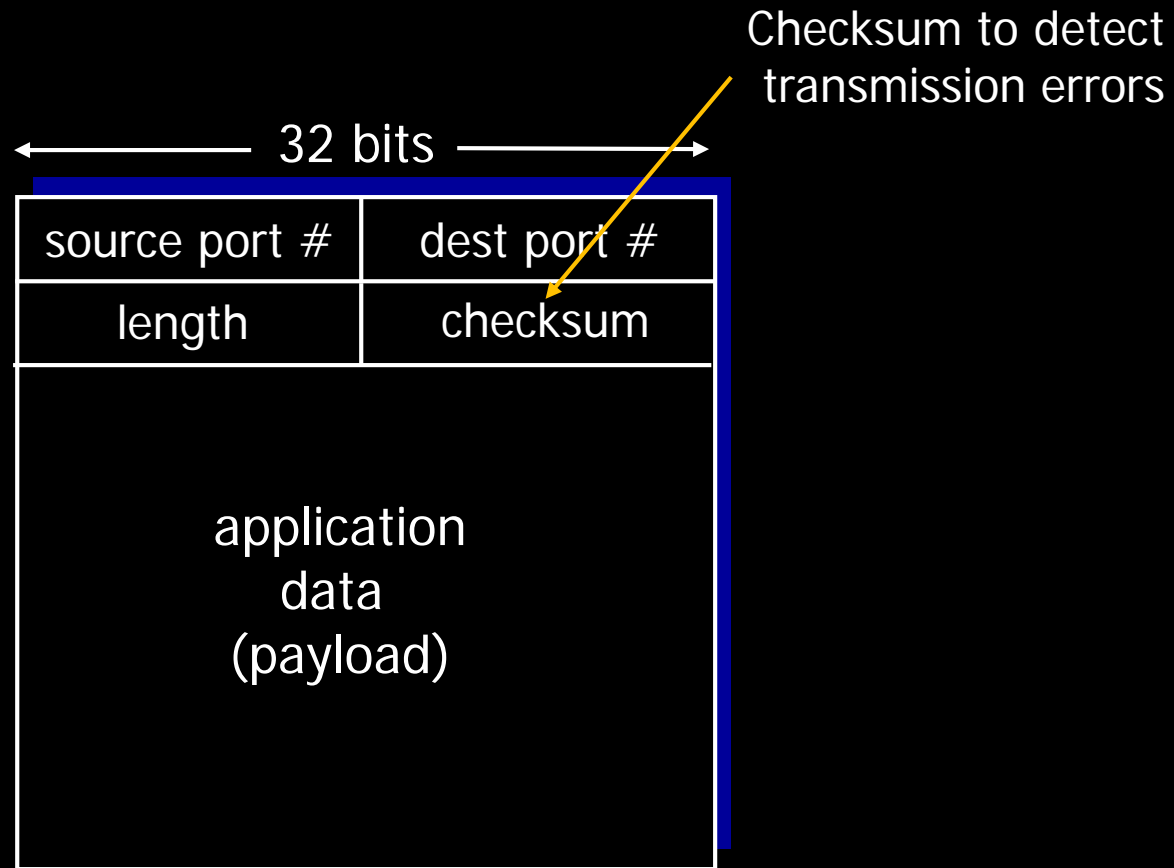
- “No frills,” “bare bones” Internet transport protocol
- Connectionless service:
 - each UDP segment handled independently of others
 - **Unreliable:** UDP segments may be lost or delivered out-of-order to app
- UDP use:
 - streaming multimedia apps (loss tolerant, rate sensitive)
 - DNS
 - SNMP
- Reliable transfer over UDP:
 - add reliability at application layer
 - application-specific error recovery!



UDP segment format

■ Advantages of UDP

- no connection establishment (which can add delay)
- simple: no connection state at sender, receiver
- small header size
- no congestion control: UDP can blast away as fast as desired



UDP segment format

■ Sender

- make a 16-bit integer checksum code of segment contents including header
- put the checksum into the checksum field and transmit the segment

■ Receiver

- compute checksum of received segment
- check if the computed checksum equals all 1's or not

Ex) splits an entire segment into 16-bit words
and add them one by one

		1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
		1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																	
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1	1
<hr/>																	
1s complement	sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum		0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

Note: when adding numbers, a carryout from the most significant bit needs to be added to the result

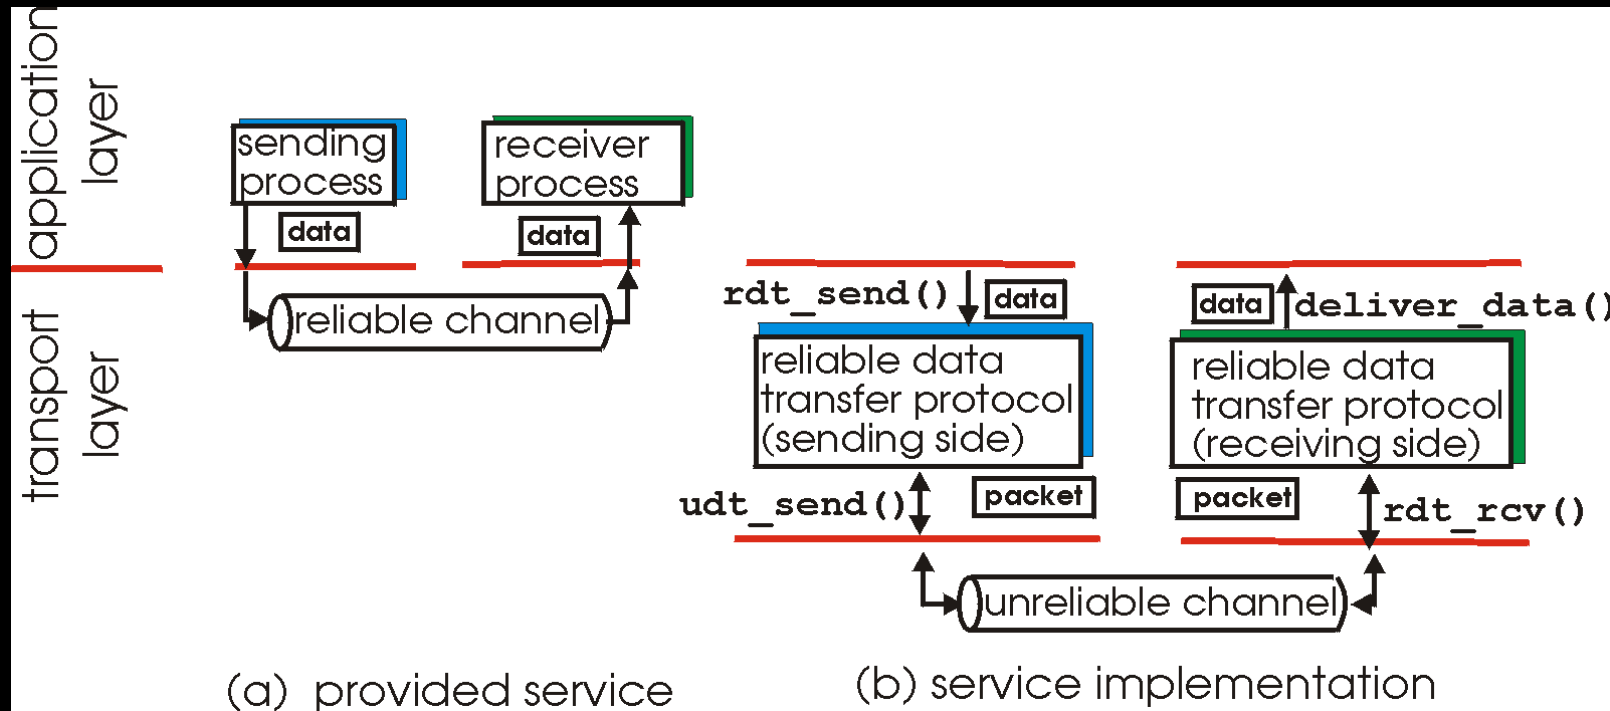
- Why checksum in UDP?



04. Reliable Data Transfer Principles

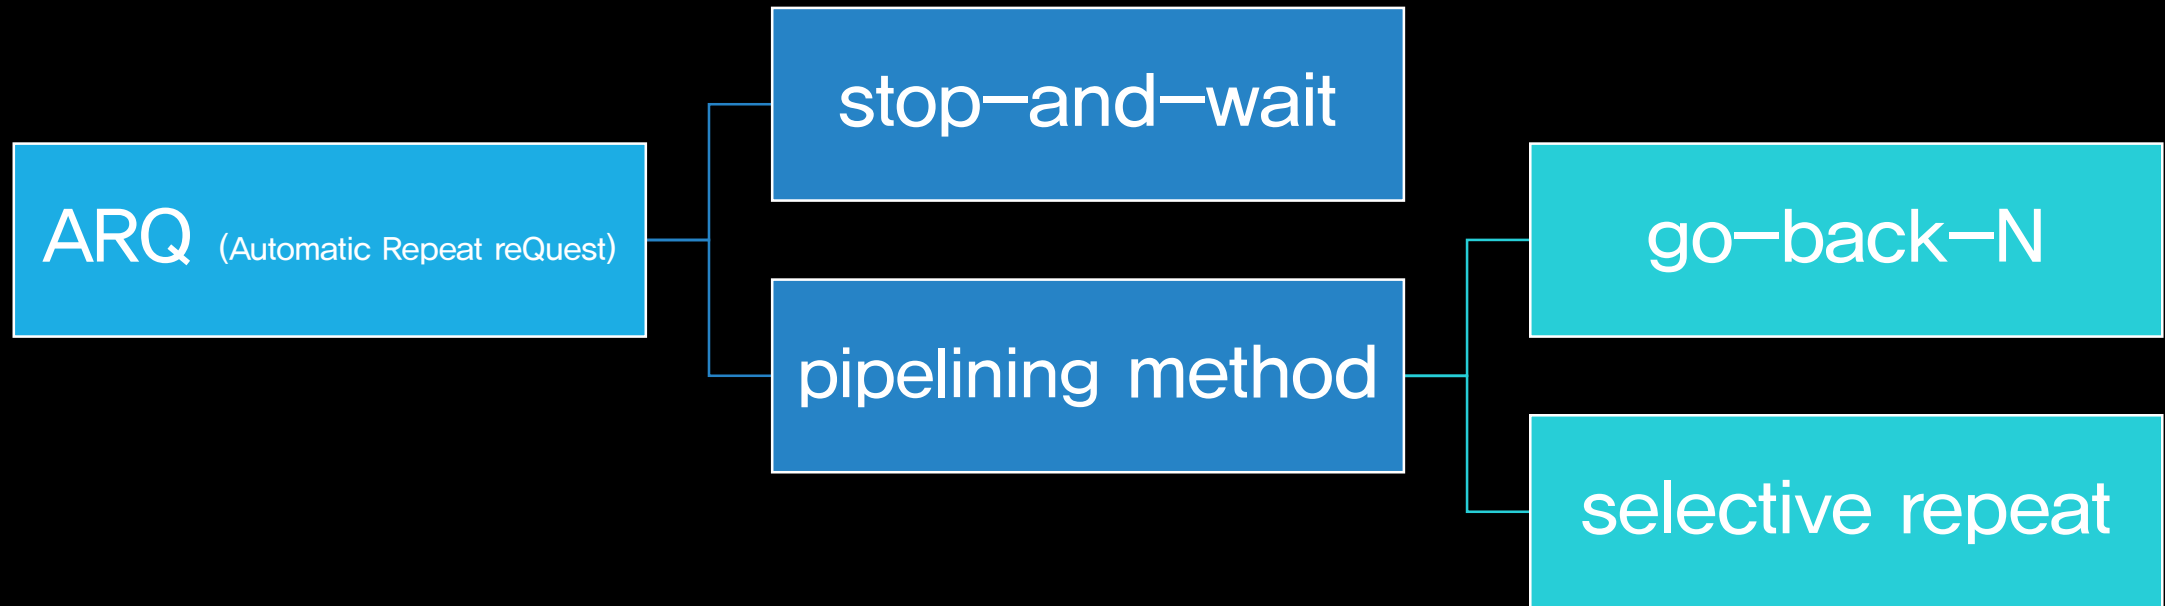
- Services abstraction (provided to the upper-layer) through a reliable channel
 - no corruption and no lost of data
 - delivered in the order in which they were sent

} Service model of TCP



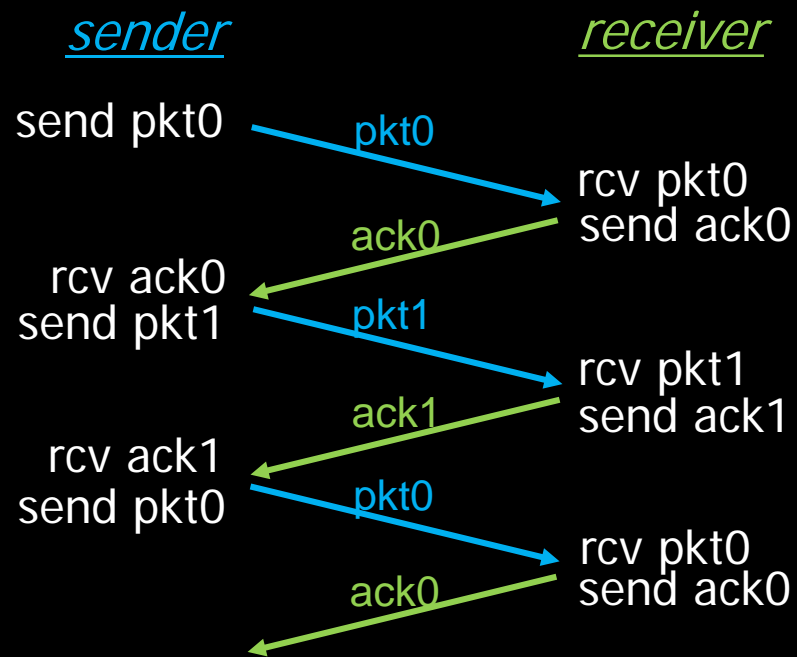
Error Types	Solution
Bit error	<ul style="list-style-type: none">• Checksum in every segment• ACK returned for successfully received packet
Packet loss (Data or ACK)	<ul style="list-style-type: none">• Timeout of sender's timer• Packet retransmission• Packet sequence number for<ol style="list-style-type: none">1) ordered delivery2) data duplication prevention

- Packet retransmission method

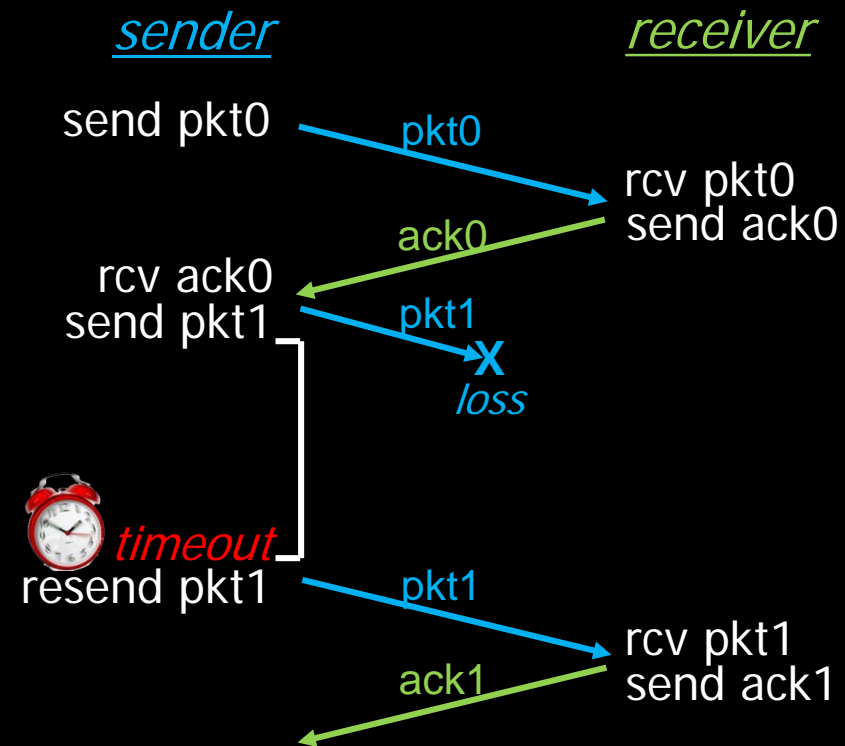


Stop-and-wait

- sender sends one packet, then waits for receiver response
- after receiving ACK, sender resumes transmission
- if timer expires without receiving ACK, sender retransmits the previous packet



(a) no loss



(b) packet loss




(c) ACK loss



(d) premature timeout/ delayed ACK

- E.g., 1 Gbps link, 15 ms propagation delay, 1kB (= 8000 bit) packet


$$D_{trans} = \frac{L}{R} = \frac{8000 \text{ bits}}{10^9 \text{ bits/sec}} = 8 \mu s$$



- U : utilization – fraction of time sender busy sending

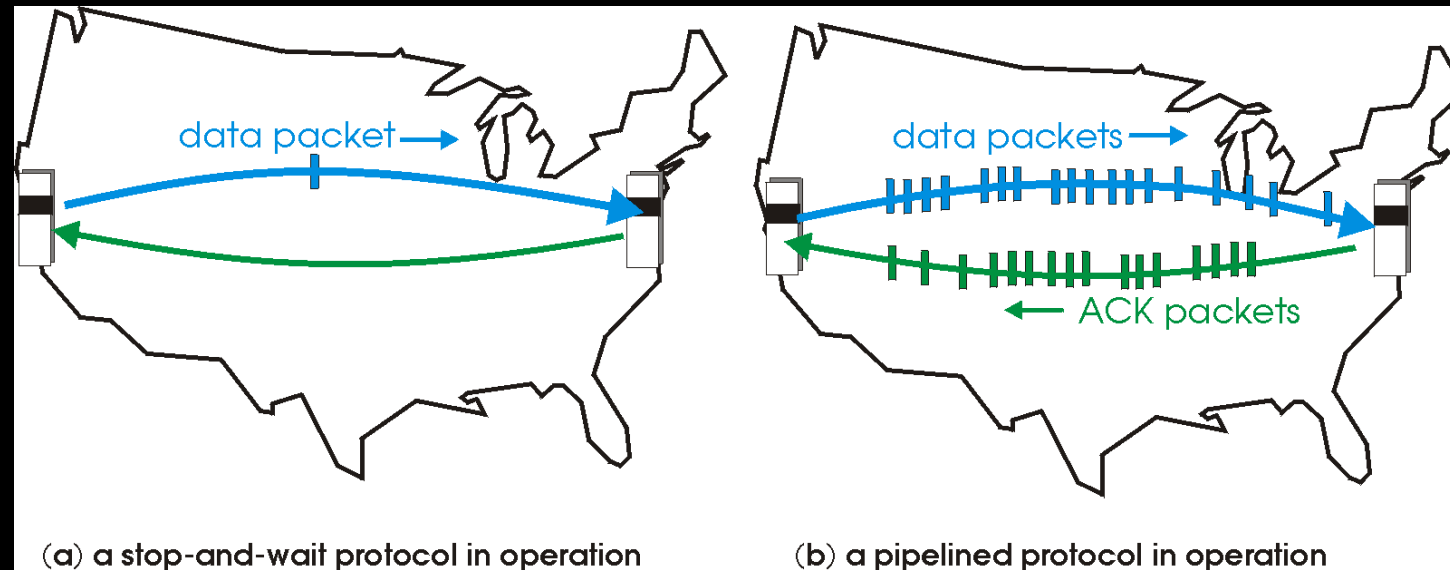
$$U = \frac{L/R}{RTT + L/R} = \frac{.008}{30.008} = 0.00027$$

- if $RTT=30ms$, 1kB packet every 30 ms: 33kB/sec throughput over 1 Gbps link
- Network protocol limits use of physical resources!

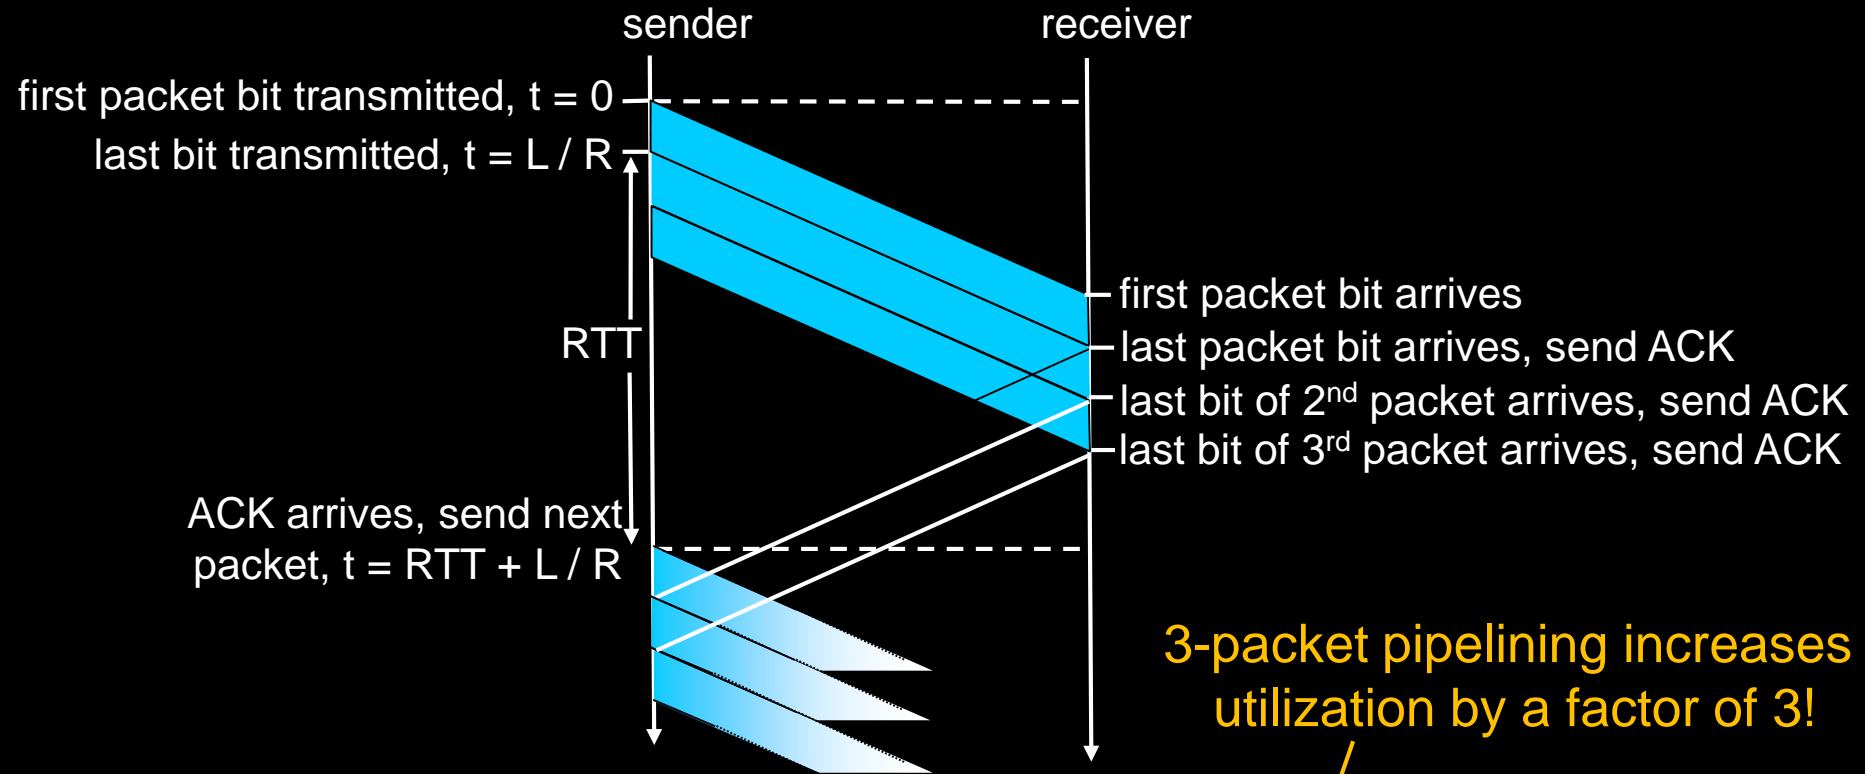


$$U = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

- **Pipelining:** sender allows multiple, “in-flight”, yet-to-be-acknowledged packets
 - range of sequence numbers must be increased
 - buffering at sender and/or receiver



- Two generic forms of pipelined protocols: *go-back-N*, *selective repeat*



$$U = \frac{3L / R}{RTT + L / R} = \frac{.024}{30.008} = 0.00081$$


Go-back-N

- Sender can have up to N unacked packets in pipeline
- Receiver only sends **cumulative ack**
 - in case a gap in seq. #, resends the last ack
- Sender has timer for oldest unacked packet
 - when timer expires, retransmit all unacked packets

Selective repeat

- Sender can have up to N unacked packets in pipeline
- Receiver sends **individual ack** for each packet
- Sender has timer for each unacked packet
 - when timer expires, retransmit only that unacked packet

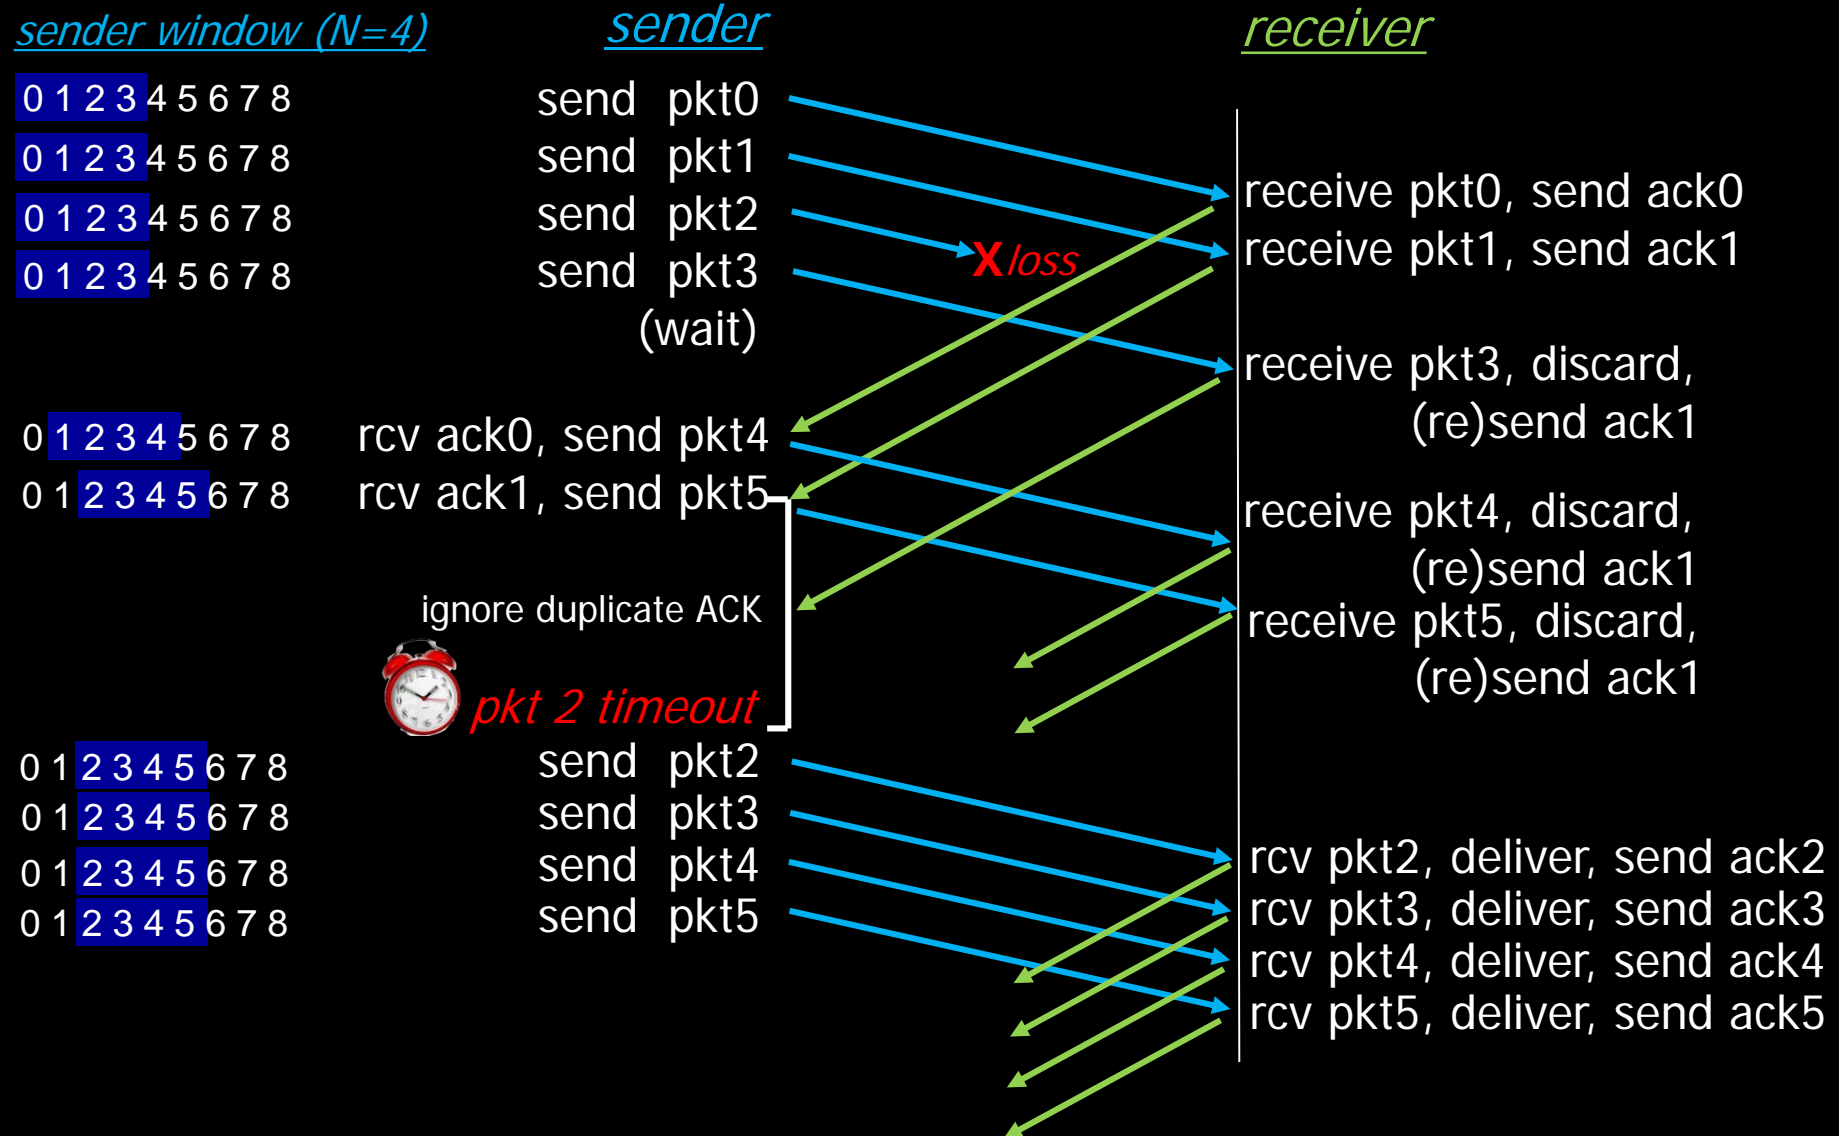
- Sequence number in packet header
- “Window” of up to N , consecutive unacked packets allowed



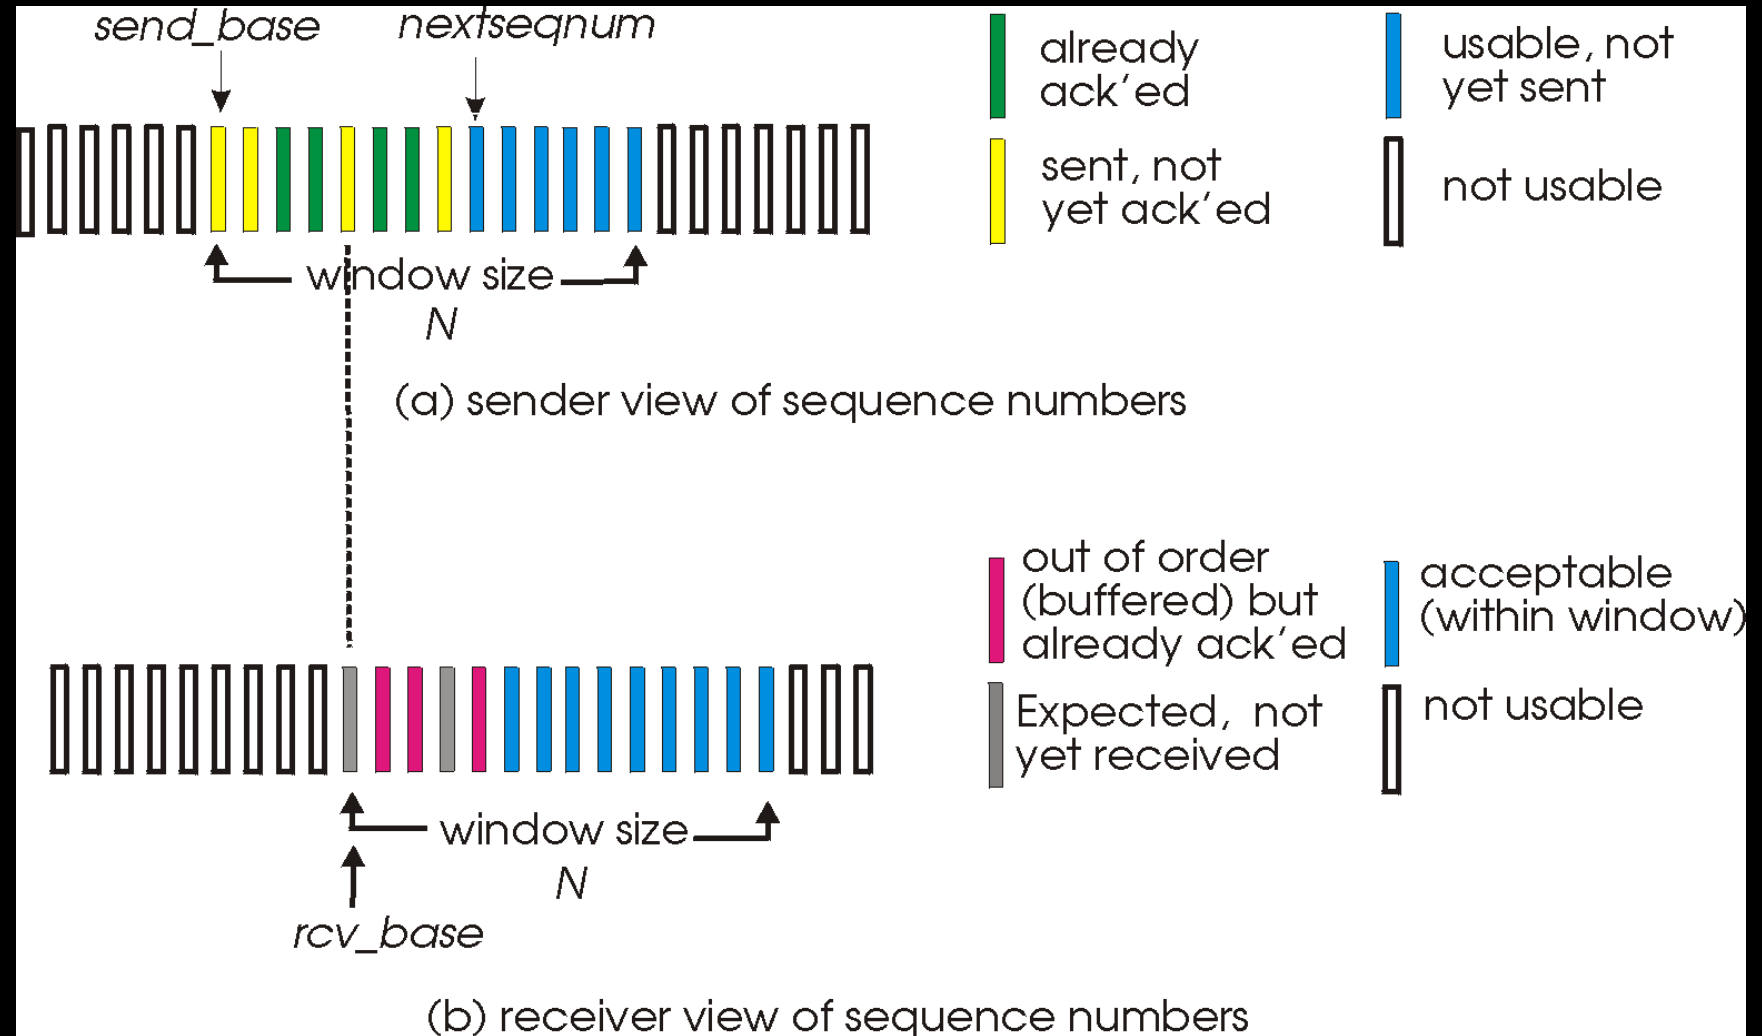
Sliding window

Window size N is dynamically changed in order to maximize the throughput, on condition that it does not cause network congestion or receiver buffer overflow

- **Cumulative ACK:** acknowledges all packets up to
- Timer for oldest in-flight packet, *send_base*
 - when timer expires for any packet with seq. # n , packet n and all higher seq. # packets in window are retransmitted



- Sender only resends packets for which ACK not received
 - sender timer for each unACKed packet
- Receiver individually acknowledges all correctly received packets
 - buffers packets, as needed, for eventual in-order delivery to upper layer



Selective Repeat Operation



sender window (N=4)

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

sender

send pkt0

send pkt1

send pkt2

send pkt3

(wait)

rcv ack0, send pkt4

rcv ack1, send pkt5

record ack3 arrived



pkt 2 timeout

send pkt2

record ack4 arrived

record ack5 arrived

receiver

receive pkt0, send ack0

receive pkt1, send ack1

receive pkt3, buffer,
send ack3

receive pkt4, buffer,
send ack4

receive pkt5, buffer,
send ack5

rcv pkt2; deliver pkt2,
pkt3, pkt4, pkt5; send ack2


X loss

Q: what happens when ack2 arrives?


example:

- Seq. #'s: 0, 1, 2, 3
- Window size=3
- Receiver sees no difference in two scenarios!
- Duplicate data accepted as new in (b)

Q: What relationship between sequence number size and window size to avoid problem in (b)?



receiver can't see sender side.
 receiver behavior identical in both cases!
something's (very) wrong!





05. Transmission Control Protocol

- Point-to-point: one sender, one receiver
- Connection-oriented service
 - Reliable transfer, in-order delivery
 - handshaking initializes sender and receiver state before data exchange
- **Pipelined transmission**
 - window size is set by TCP congestion and flow control
 - **congestion control**: transmission rate controlled not to make congestion in network
 - **flow control**: sender will not overwhelm receiver
 - MSS: maximum segment size
- **Full-duplex connection**
 - bi-directional data flow in same connection

TCP Segment Header Format

Bit #	0	7	8	15	16	23	24	31
0	Source Port				Destination Port			
32	Sequence Number							
64	Acknowledgment Number							
96	Data Offset	Res	Flags		Window Size			
128	Header and Data Checksum				Urgent Pointer (not used)			
160...	Options							

U

A

P

R

S

F

Reserved for future use

Data Offset = Header length

Internet checksum (as in UDP)



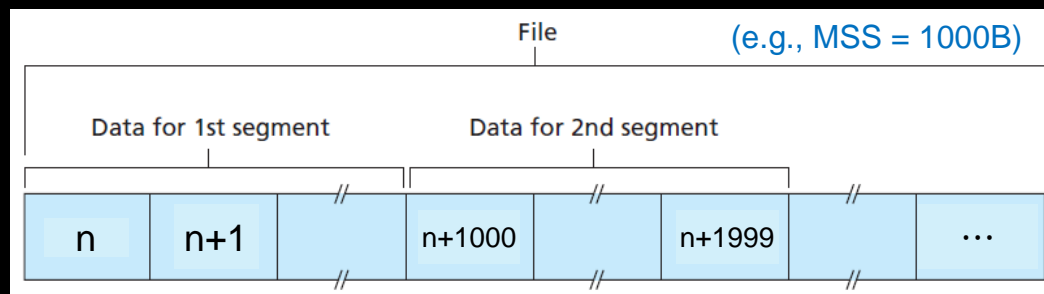
cf.

UDP Datagram Header Format

Bit #	0	7	8	15	16	23	24	31
0	Source Port							
32	Destination Port							
	Length				Header and Data Checksum			

Sequence number

- byte stream “number” of first byte in segment’s data




Acknowledgment number

- sequence number of the next segment expected by receiver
- cumulative ACK

outgoing segment from sender

source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer



sent
ACKed


sent, not-
yet ACKed
("in-
flight")


usable
but not
yet sent

not
usable

incoming segment to sender

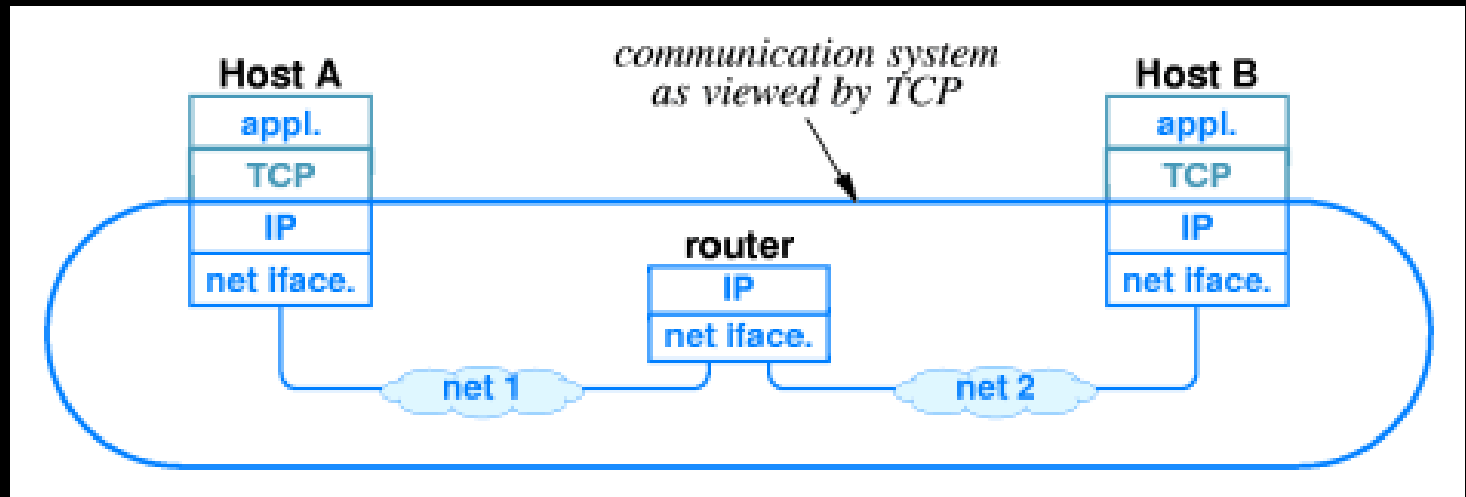
source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer





simple telnet scenario

- TCP creates reliable data transfer service on top of IP's unreliable service
 - pipelined segments
 - cumulative acks
 - single retransmission timer



출처 -

https://www.google.co.kr/url?sa=i&rct=j&q=&esrc=s&source=images&cd=&cad=rja&uact=8&ved=2ahUKEwjf9czB0PXbAhUI5rwKHxbIBqkQjRx6B-AgBEAU&url=https%3A%2F%2Fwww.cs.csustan.edu%2F~john%2Fclasses%2Fprevious_semesters%2FCS3000_Communication_Networks%2F2007_02_Spring%2FNotes%2Fchap25.html&psig=AOvVaw2QofBlqkTFxG8_J4eyPGI&ust=1530250010414409

■ TCP sender

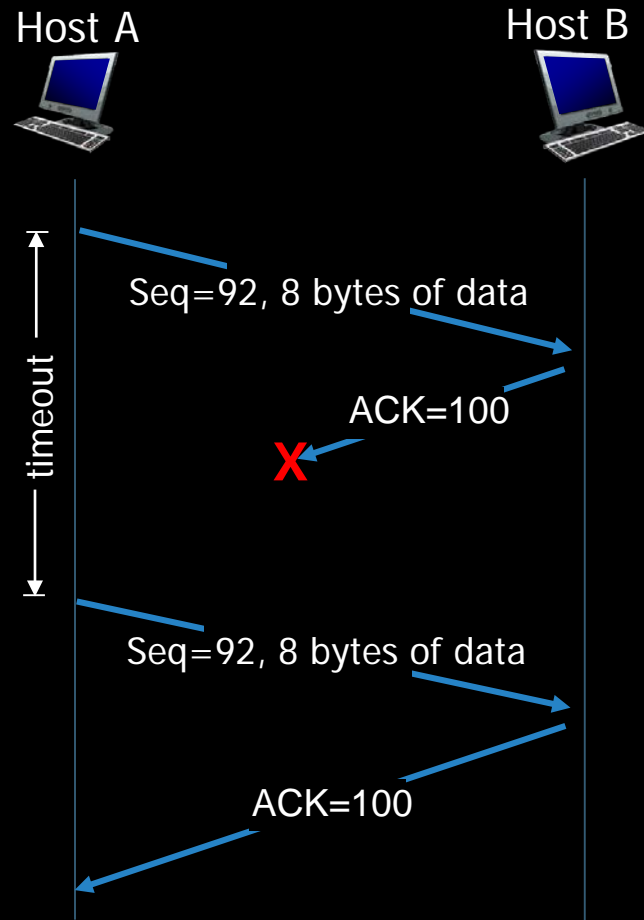
- when receiving data from app
 - create segment with seq. number
 - send it and start timer if not already running
- when timer expires
 - retransmit the segment that caused timeout and restart timer
- when receiving ack
 - update ACKed packet list
 - start timer if there are still unacked segments

■ TCP receiver

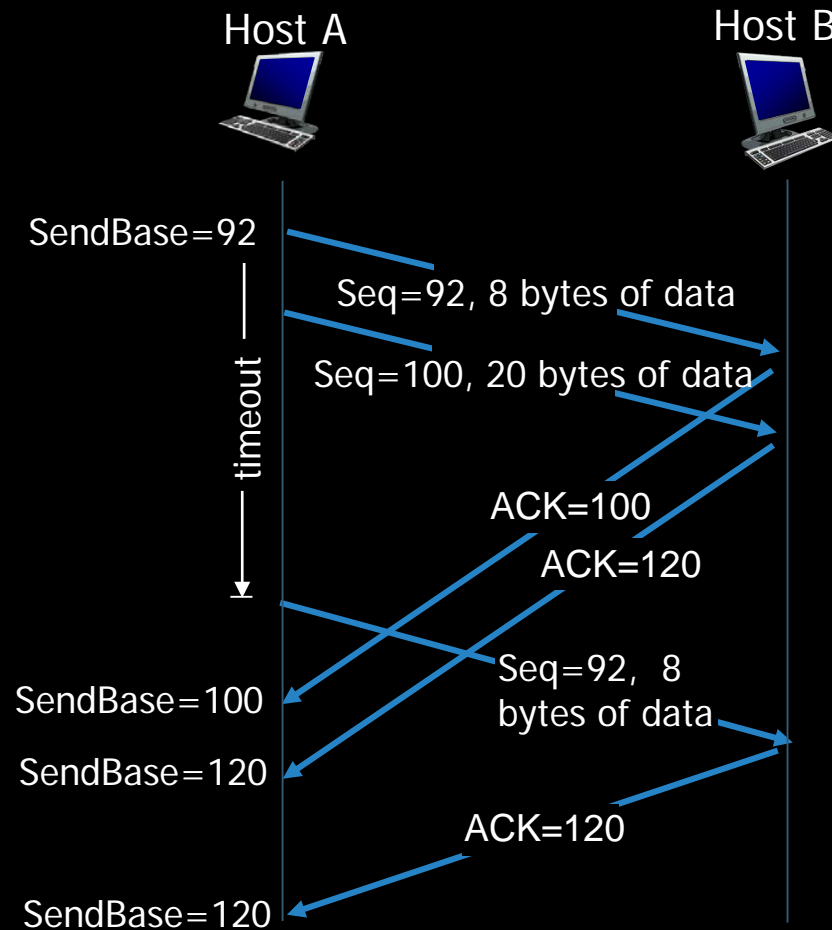
- when receiving data from sender
 - if discover bit error through checksum, then drop the packet
 - checking seq. # and there being no gap, send cumulative ACK
 - if packet is duplicated, drop it

■ How receiver handles out-of-order segments?


- TCP specification does not say explicitly
 - up to implementer



lost ACK scenario



premature timeout



How to set TCP timeout value?

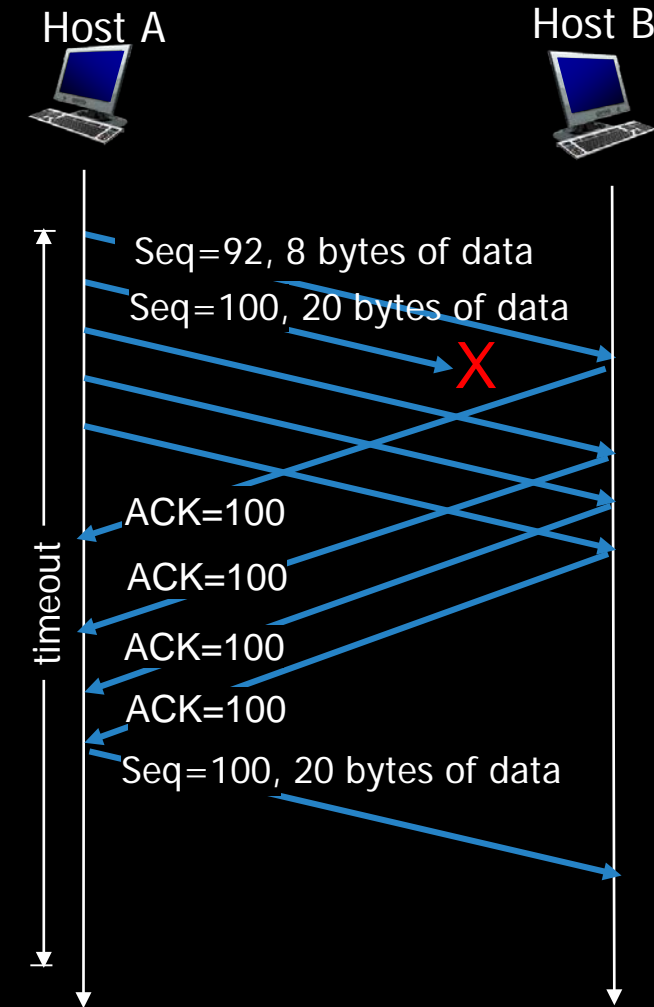
- must longer than RTT
- `SampleRTT`
 - measured time from segment transmission until ACK receipt
- `EstimatedRTT`
 - `SampleRTT` varies
 - moving average several recent measurements

$$\text{EstimatedRTT} = (1-\alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$



$$\text{TimeoutInterval} = \text{EstimatedRTT} + \text{"safety margin"}$$

- Timeout period often relatively long:
 - long delay before resending lost packet
- Another way to detect segment lost
 - whenever a segment is received, receiver sends ACK to request the segment that has the next sequence number of the last in-order segment
 - sender often sends many segments back-to-back
 - if a segment is lost, there will likely be many duplicate ACKs

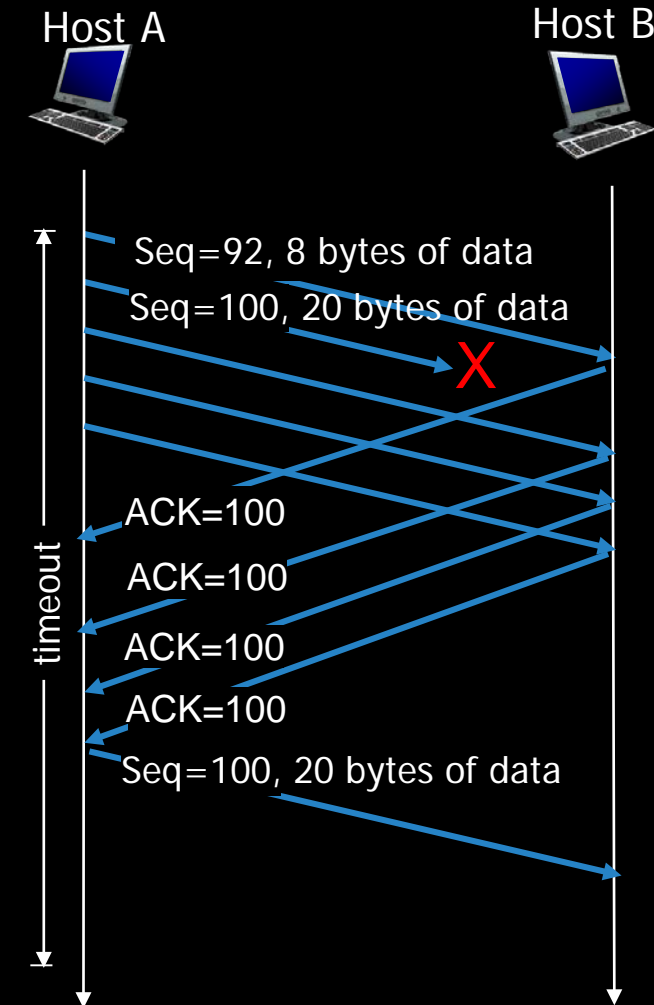


- **TCP fast retransmit**

- if sender receives 3 duplicate ACKs (except the first normal ACK) for same data, resend unacked segment with smallest seq. number

- As a result, TCP retransmissions triggered by:

- 1) timeout events
- 2) duplicate acks



TCP Header

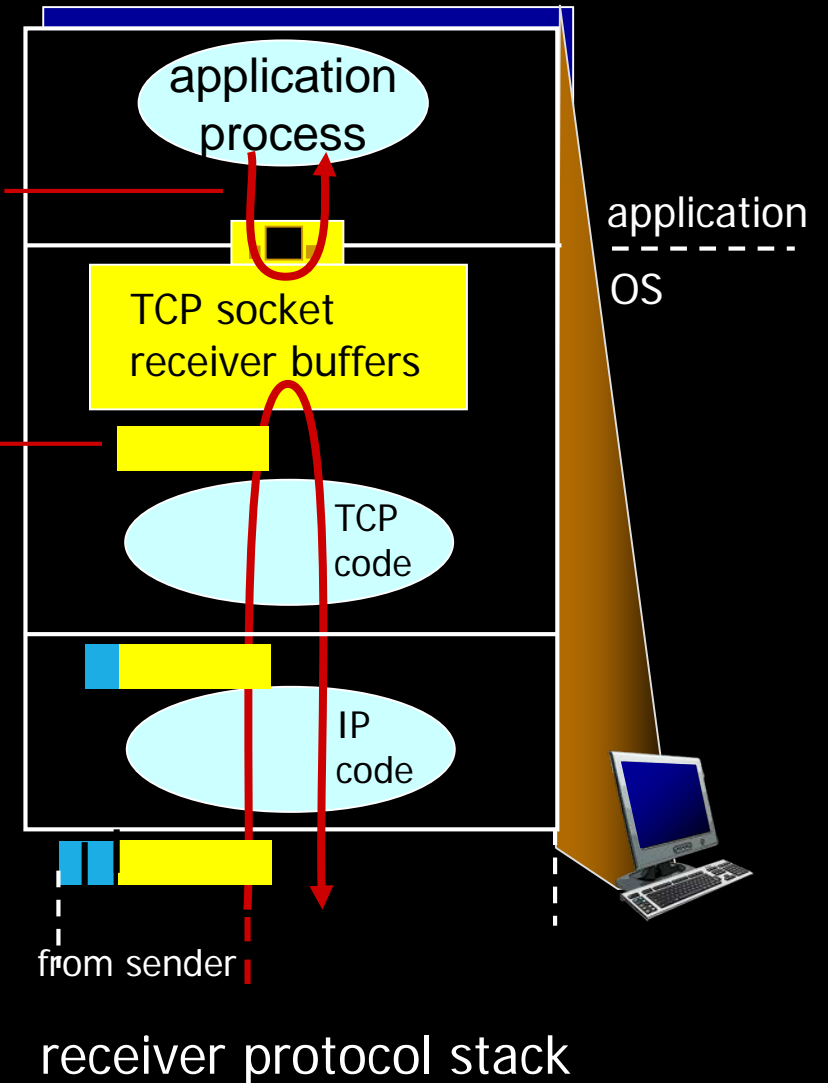
source port #		dest port #	
sequence number			
acknowledgement number			
h_len		flags	receive window
checksum		urgent pointer	

flow control

Receiver controls sender, so sender won't overflow receiver's buffer by transmitting too much, too fast

application may remove data from TCP socket buffers

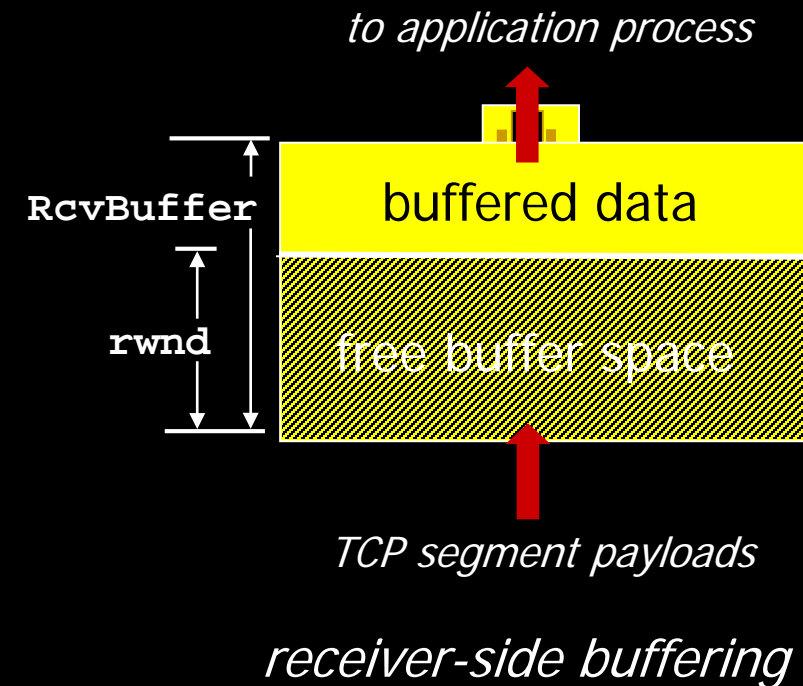
... slower than TCP receiver is delivering (sender is sending)



- Receiver advertises free buffer space by including **rwnd** value in TCP header of receiver-to-sender segments
 - RcvBuffer** size set via socket options (typical default is 4096 bytes)
 - many OSs autoadjust **RcvBuffer**
- Sender limits amount of unacked (“in-flight”) data to receiver’s **rwnd** value
- Guarantees receive buffer will not overflow

TCP Header


source port #		dest port #	
sequence number			
acknowledgement number			
h_len		flags	receive window
checksum		urgent pointer	



- Before exchanging data, sender and receiver
 - agree to establish connection
 - agree on connection parameters

TCP Header

source port #		dest port #	
sequence number			
acknowledgement number			
h_len		flags	receive window
checksum		urgent pointer	

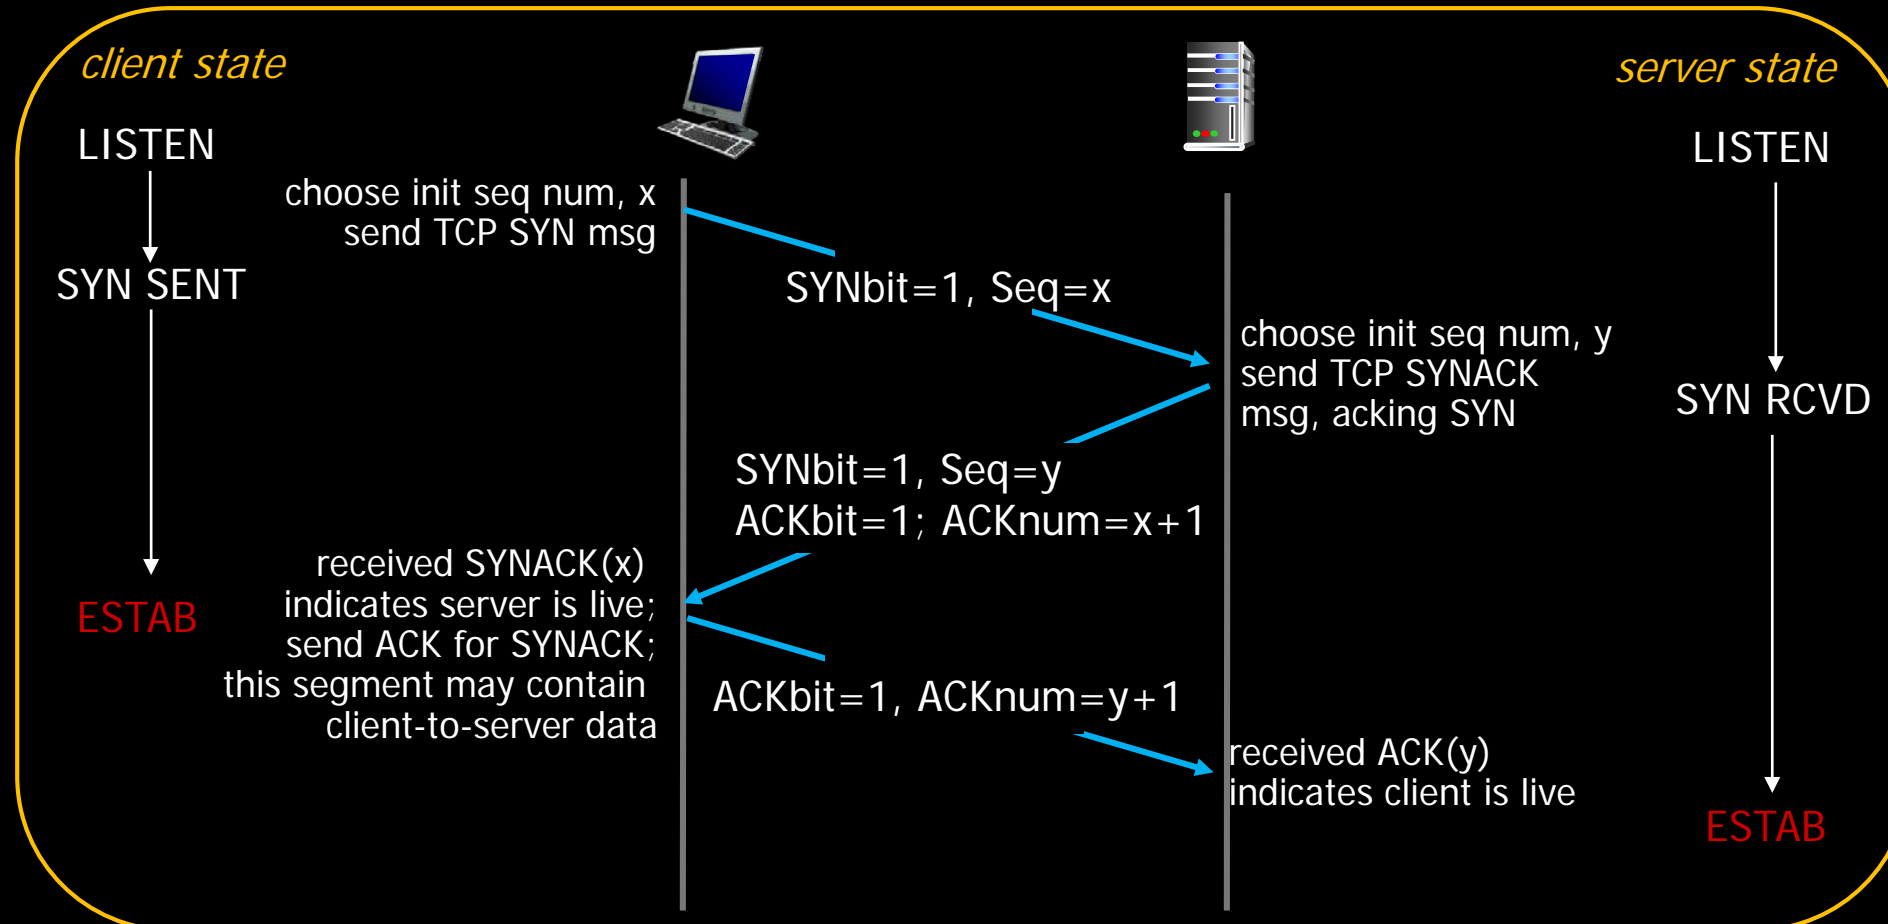


- URG: urgent data (generally not used)
- ACK: ACK # valid
- PSH: push data now (generally not used)

Three-way handshake

TCP Header

source port #	dest port #
sequence number	
acknowledgement number	
h_len	flags
checksum	urgent pointer



- Client and server each close their side of connection

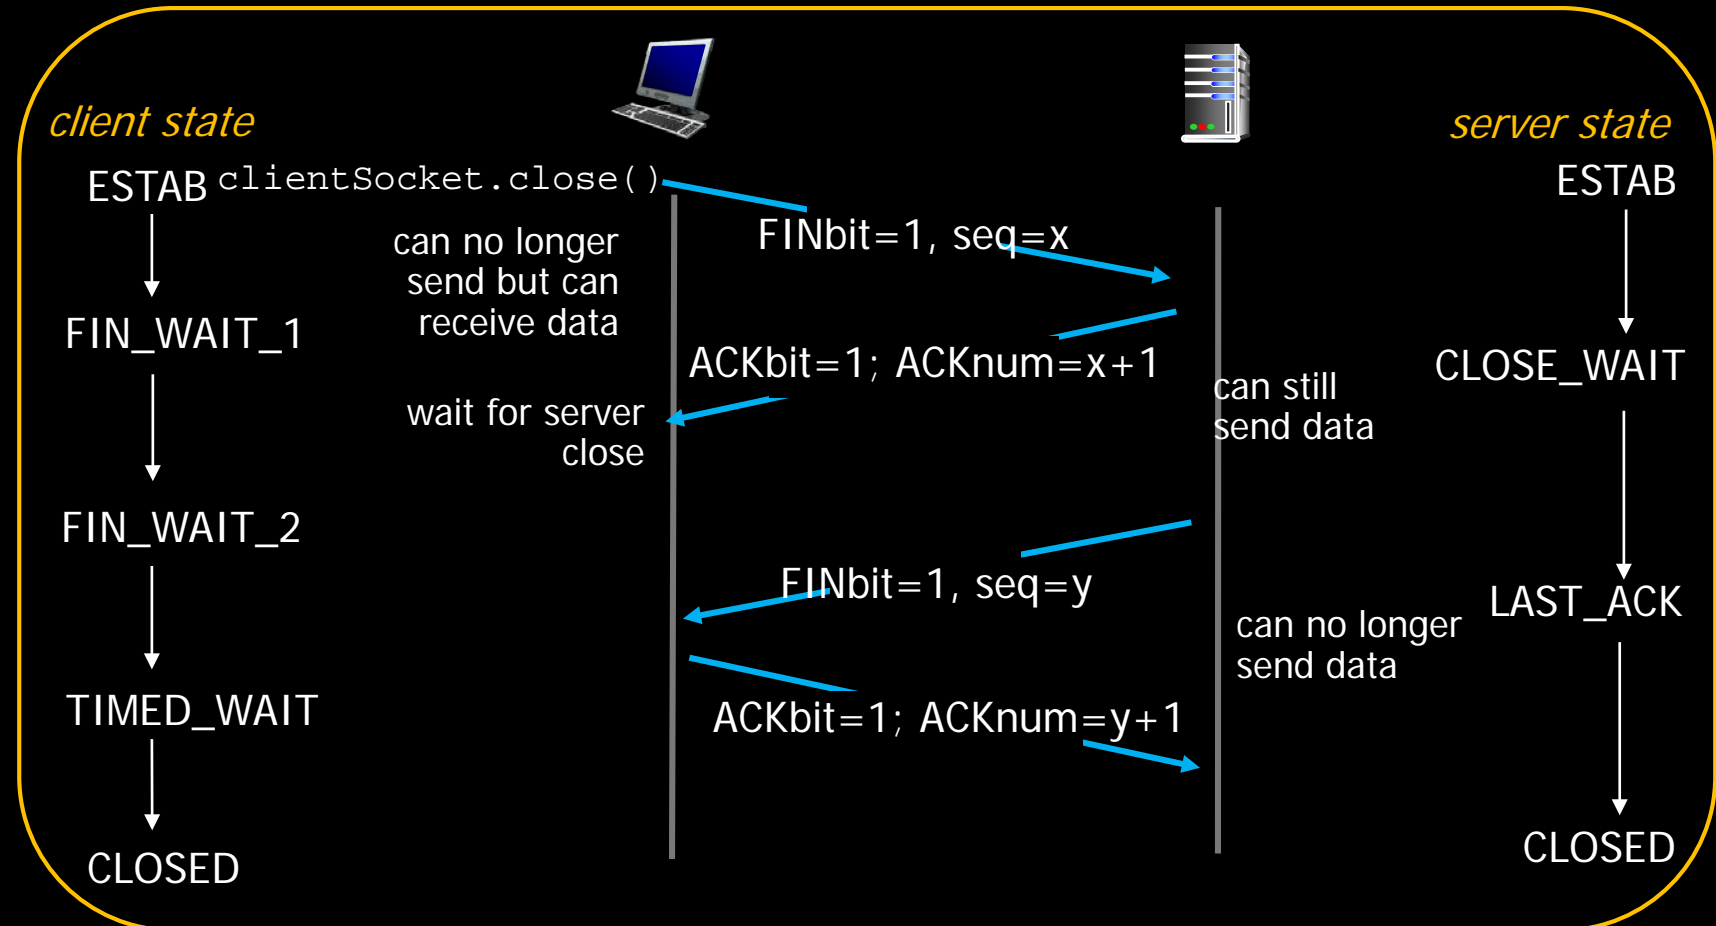
- send TCP segment with FIN bit = 1

- Respond to received FIN with ACK

- on receiving FIN, ACK can be combined with own FIN

TCP Header

source port #	dest port #
sequence number	
acknowledgement number	
h_len	flags
receive window	
checksum	urgent pointer






06. Congestion Control

- When the offered load in an distributed sharing system exceeds total system capacity, the effective load will go to zero (collapses) as load increases

➡ [Congestion collapse](#) (1980)

Congestion

- “Too many sources sending too much data too fast for *network* to handle”
- Different from flow control!
 - congestion control is global issue
 - flow control is point-to-point issue
- Manifestations:
 - lost packets (buffer overflow at routers)
 - long delays (queueing in router buffers)



출처 - A. Afanasyev et al., “Host-to-Host Congestion Control for TCP,” IEEE Comm. Surveys & Tutorials, Vol. 12, No. 3, 3rd Quarter, 2010, pp. 304–342

- Two approaches towards congestion control:

end-to-end


- No explicit feedback from network
- Congestion inferred from end-system observed loss, delay
- Approach taken by TCP

network-assisted

- Routers provide feedback to end systems
 - single bit indicating congestion (SNA, DECbit, TCP/IP ECN, ATM)
 - explicit rate for sender to send at

- **AIMD** approach: sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs
 - **additive increase**: increase **cwnd** by 1 MSS every RTT until loss detected
 - **multiplicative decrease**: cut **cwnd** in half after loss

AIMD saw tooth
behavior: probing
for bandwidth




- Congestion window (`cwnd`)


- the rate a TCP sender can send traffic (the amount of packets in transit)

$$\text{LastByteSent} - \text{LastByteAcked} \leq \min\{\text{cwnd}, \text{rwnd}\}$$

- `cwnd` is dynamic,
function of perceived network congestion




- When connection begins, increase rate exponentially until first loss event:
 - initially `cwnd` = 1 MSS
 - double `cwnd` every RTT
 - done by incrementing `cwnd` for every ACK received
- **Summary**: initial rate is slow but ramps up exponentially fast



- Q: When should the exponential growth in the slow-start state end?
- A: When `cwnd` gets to 1/2 of its value before congestion was detected


- Implementation
 - variable `ssthresh`
 - on loss event, `ssthresh` is set to 1/2 of `cwnd` just before loss event
 - surpassing `ssthresh`, `cwnd` increases by just a single MSS every RTT

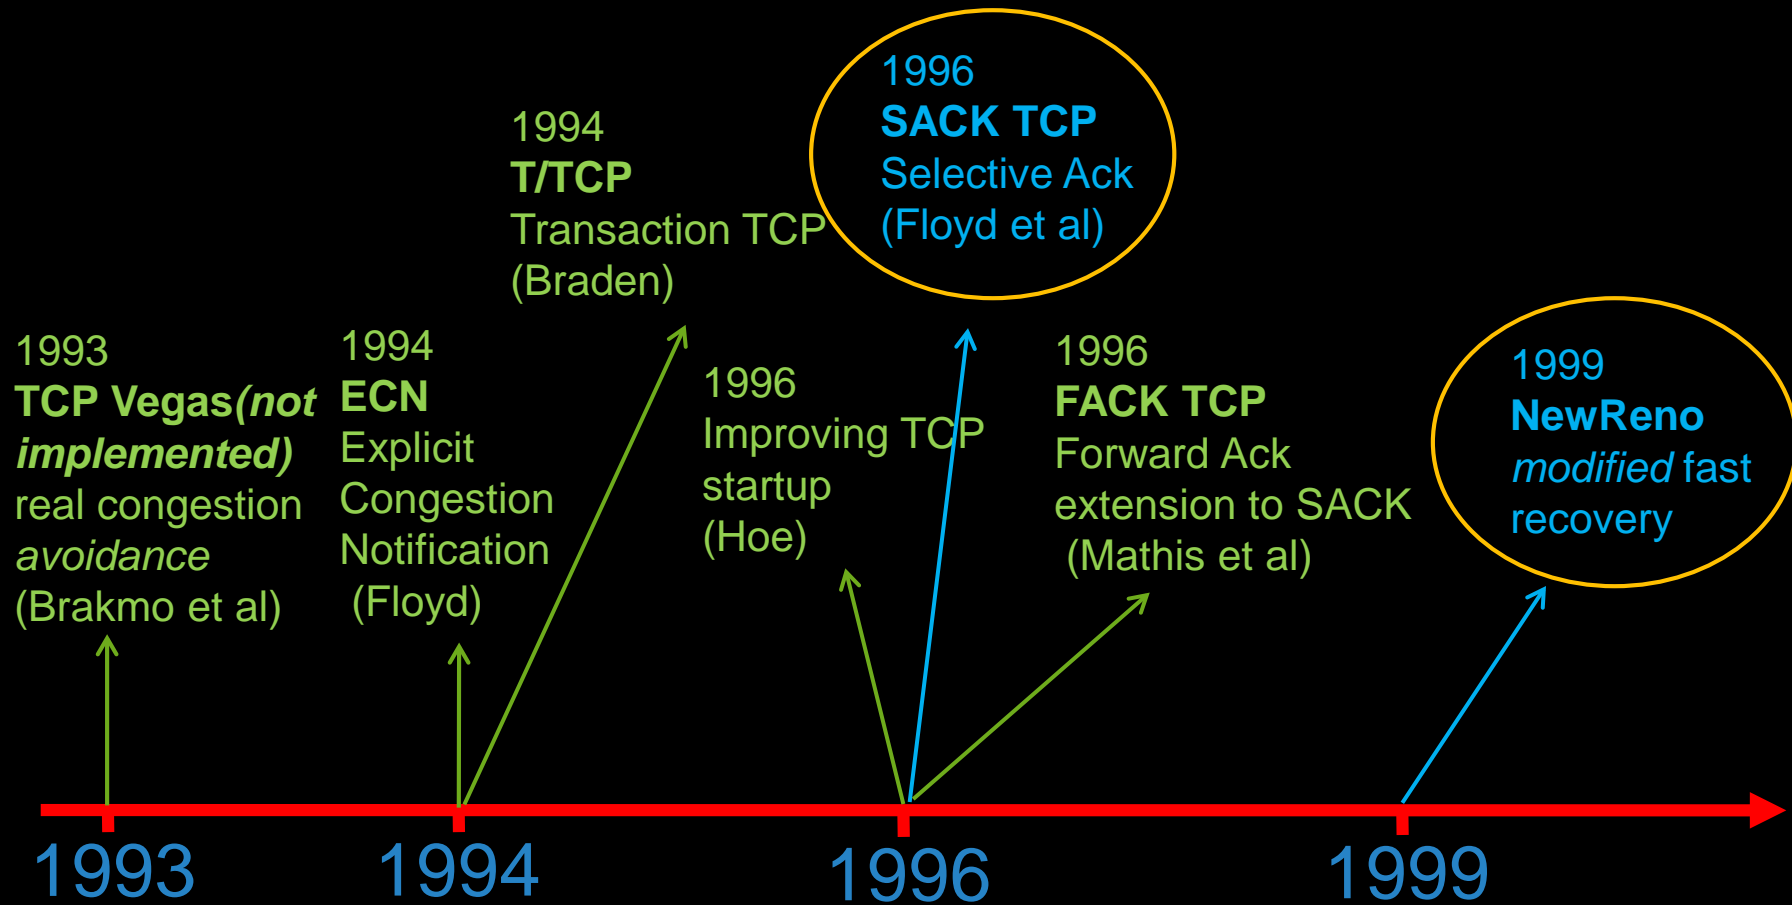
TCP Slow-Start & Congestion Avoidance





07. TCP Congestion Control Algorithm






TCP Variant	Section	Year	Base	Added/Changed Modes or Features	Mod ¹	Status	Implementation			
							BSD ²	Linux	Win	Mac
<i>TCP Tahoe</i> [14]	II-A	1988	RFC793	Slow Start, Congestion Avoidance, Fast Retransmit	S	Obsolete Standard	>4.3	1.0		
<i>TCP-DUAL</i> [15]	II-B	1992	Tahoe	Queuing delay as a supplemental congestion prediction parameter for Congestion Avoidance	S	Experimental				
<i>TCP Reno</i> [16], [17]	II-C	1990	Tahoe	Fast Recovery	S	Standard	>4.3 >F2.2	> 1.3.90	>95/NT	
<i>TCP NewReno</i> [18], [19]	II-D	1999	Reno	Fast Recovery resistant to multiple losses	S	Standard	>F4	> 2.1.36		>10.4.6 (opt)
<i>TCP SACK</i> [20]	II-E	1996	RFC793	Extended information in feedback messages	P+S+R	Standard	>S2.6, >N1.1, >F2.1R	> 2.1.90	> 98	> 10.4.6
<i>TCP FACK</i> [21]	II-F	1996	Reno, SACK	SACK-based loss recovery algorithm	S	Experimental	>N1.1	>2.1.92		
<i>TCP-Vegas</i> [22]	II-G	1995	Reno	Bottleneck buffer utilization as a primary feedback for the Congestion Avoidance and secondary for the Slow Start	S	Experimental		> 2.2.10		
<i>TCP-Vegas+</i> [23]	II-H	2000	NewReno, Vegas	Reno/Vegas Congestion Avoidance mode switching based of RTT dynamics	S	Experimental				
<i>TCP-Veno</i> [24]	II-I	2002	NewReno, Vegas	Reno-type Congestion Avoidance and Fast Recovery increase/decrease coefficient adaptation based on bottleneck buffer state estimation	S	Experimental		> 2.6.18		
<i>TCP-Vegas A</i> [25]	II-J	2005	Vegas	Adaptive bottleneck buffer state aware Congestion Avoidance	S	Experimental				

¹ TCP specification modification: S = the sender reactions, R = the receiver reactions, P = the protocol specification

² S for Sun, F for FreeBSD, N for NetBSD

- Features
 - slow start
 - congestion avoidance
 - fast retransmit
 - retransmission upon 3 duplicate ACKs
- The slow start begins (i.e., $cwnd = 1$ MSS) on
 - timeout
 - fast retransmit



출처 - A. Afanasyev et al., "Host-to-Host Congestion Control for TCP," IEEE Comm. Surveys & Tutorials, Vol. 12, No. 3, 3rd Quarter, 2010, pp. 304-342

- Limitation of Tahoe
 - no differentiates losses indicated by timeout and duplicate ACKs
 - always sets `cwnd` to 1 MSS
- Two responses to packet losses by Reno
 - loss indicated by timeout
 - `cwnd` set to 1 MSS
 - **slow start** to `ssthresh`, then congestion avoidance begins
 - loss indicated by fast retransmit
 - `cwnd` cut in half
 - **fast recovery**, then congestion avoidance


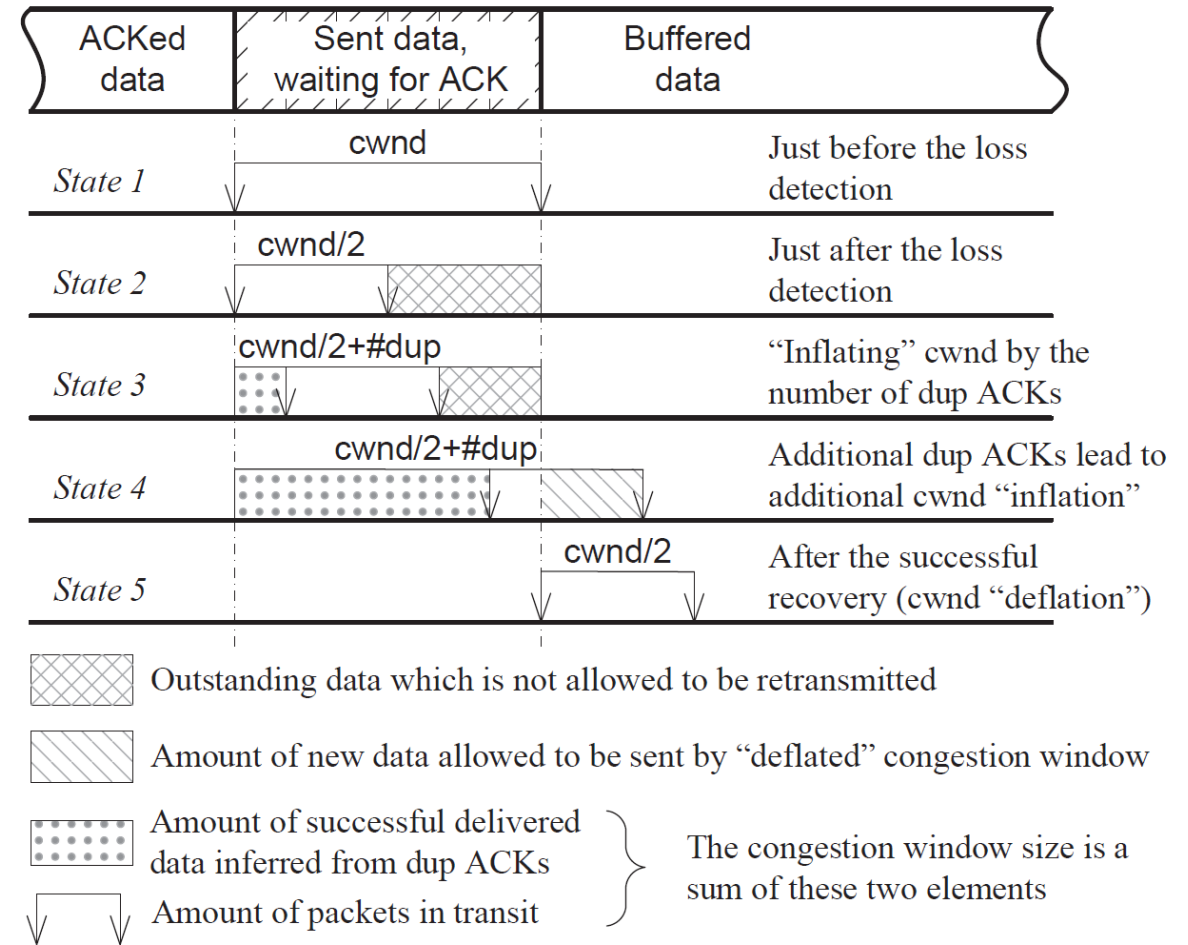



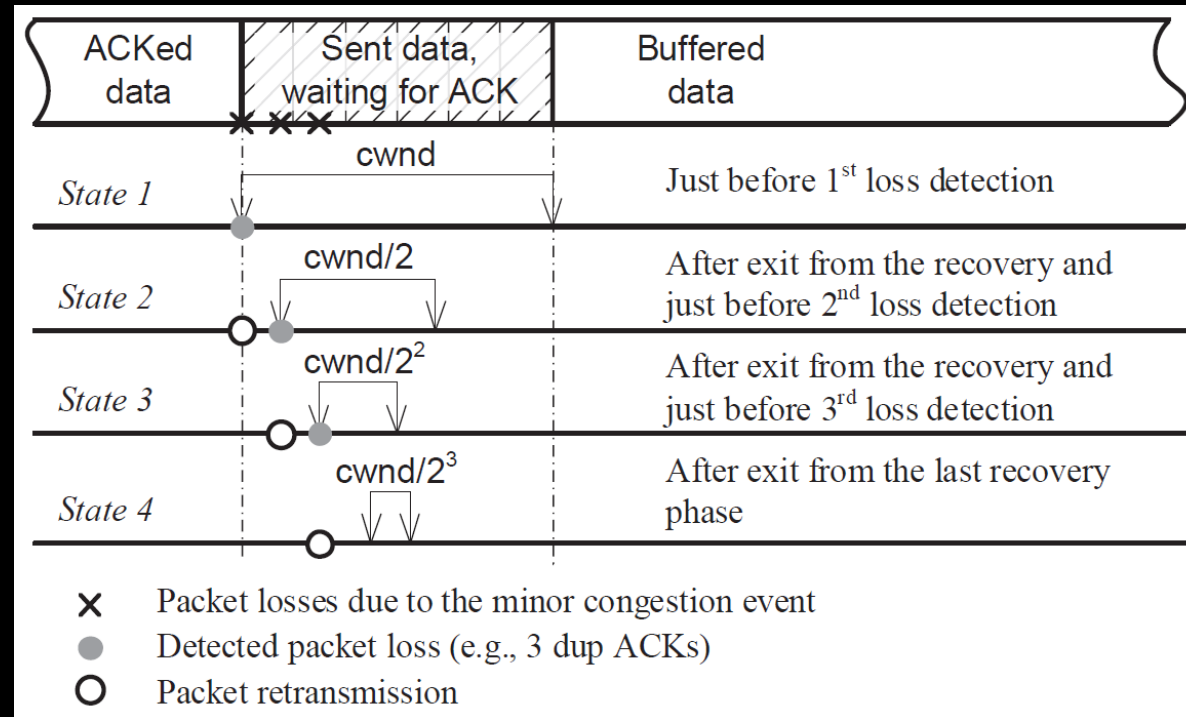
Fig. 15. Congestion window dynamics of TCP Reno (SS: the *Slow Start* phase, CA: the *Congestion Avoidance* phase, FR: the *Fast Recovery* phase)

- Half **cwnd** maintained until a non-duplicate ACK is received
 - amount of packets in transit kept from decreasing more than expected
- **Operation**
 - Upon the event of 3 duplicate ACKs,
 - after reducing **cwnd** to 1/2,
 - the algorithm not only retransmits the oldest unacknowledged packet,
 - but also inflates the congestion window by the number of duplicate packets.
 - A non-duplicate ACK makes the congestion avoidance start.



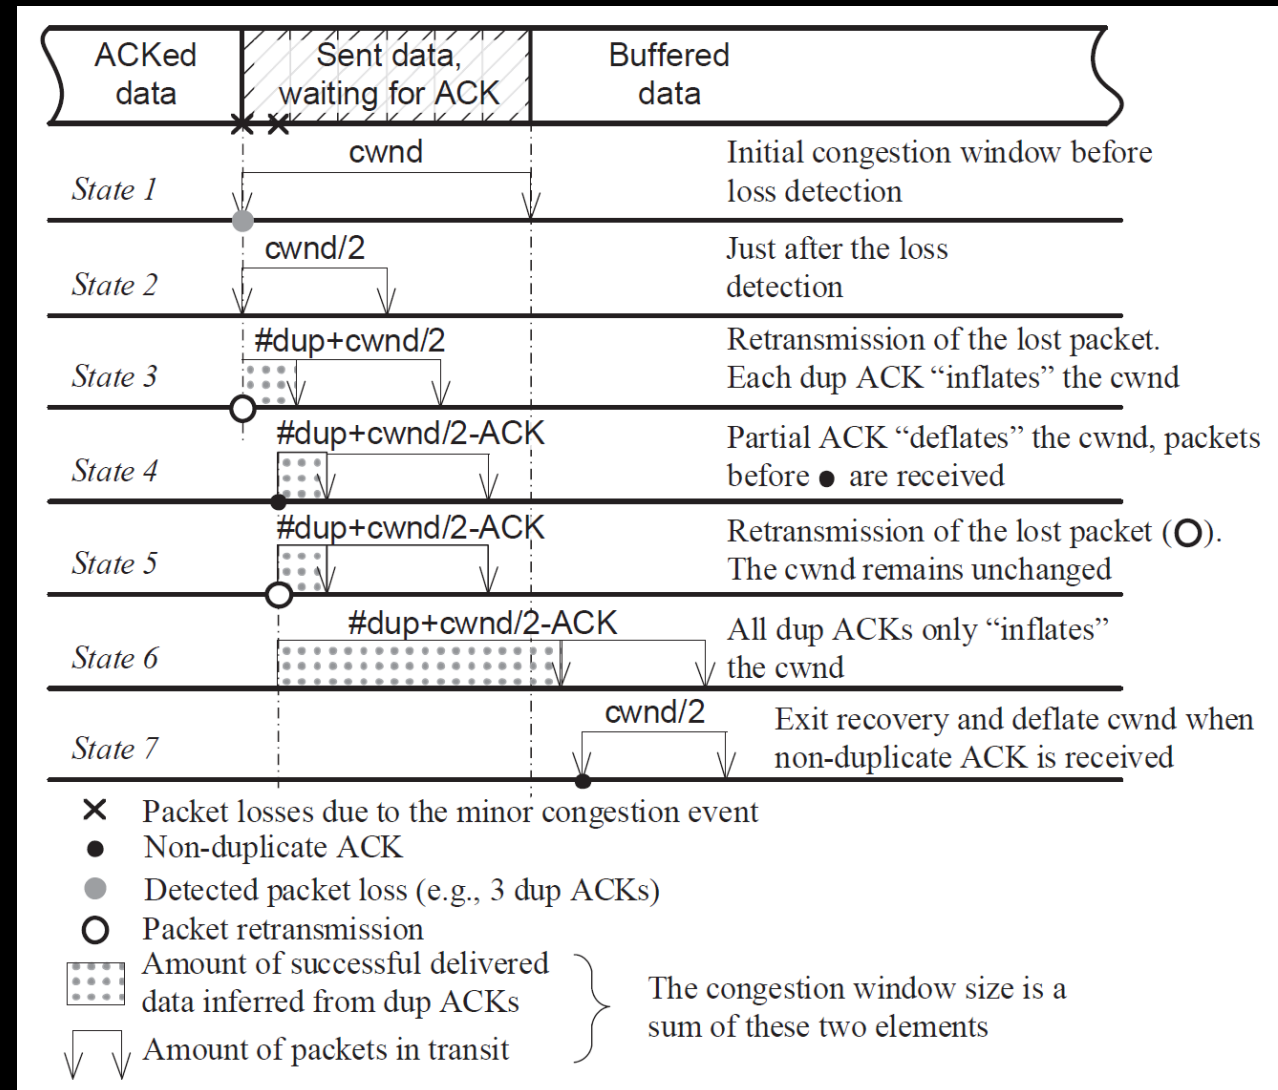


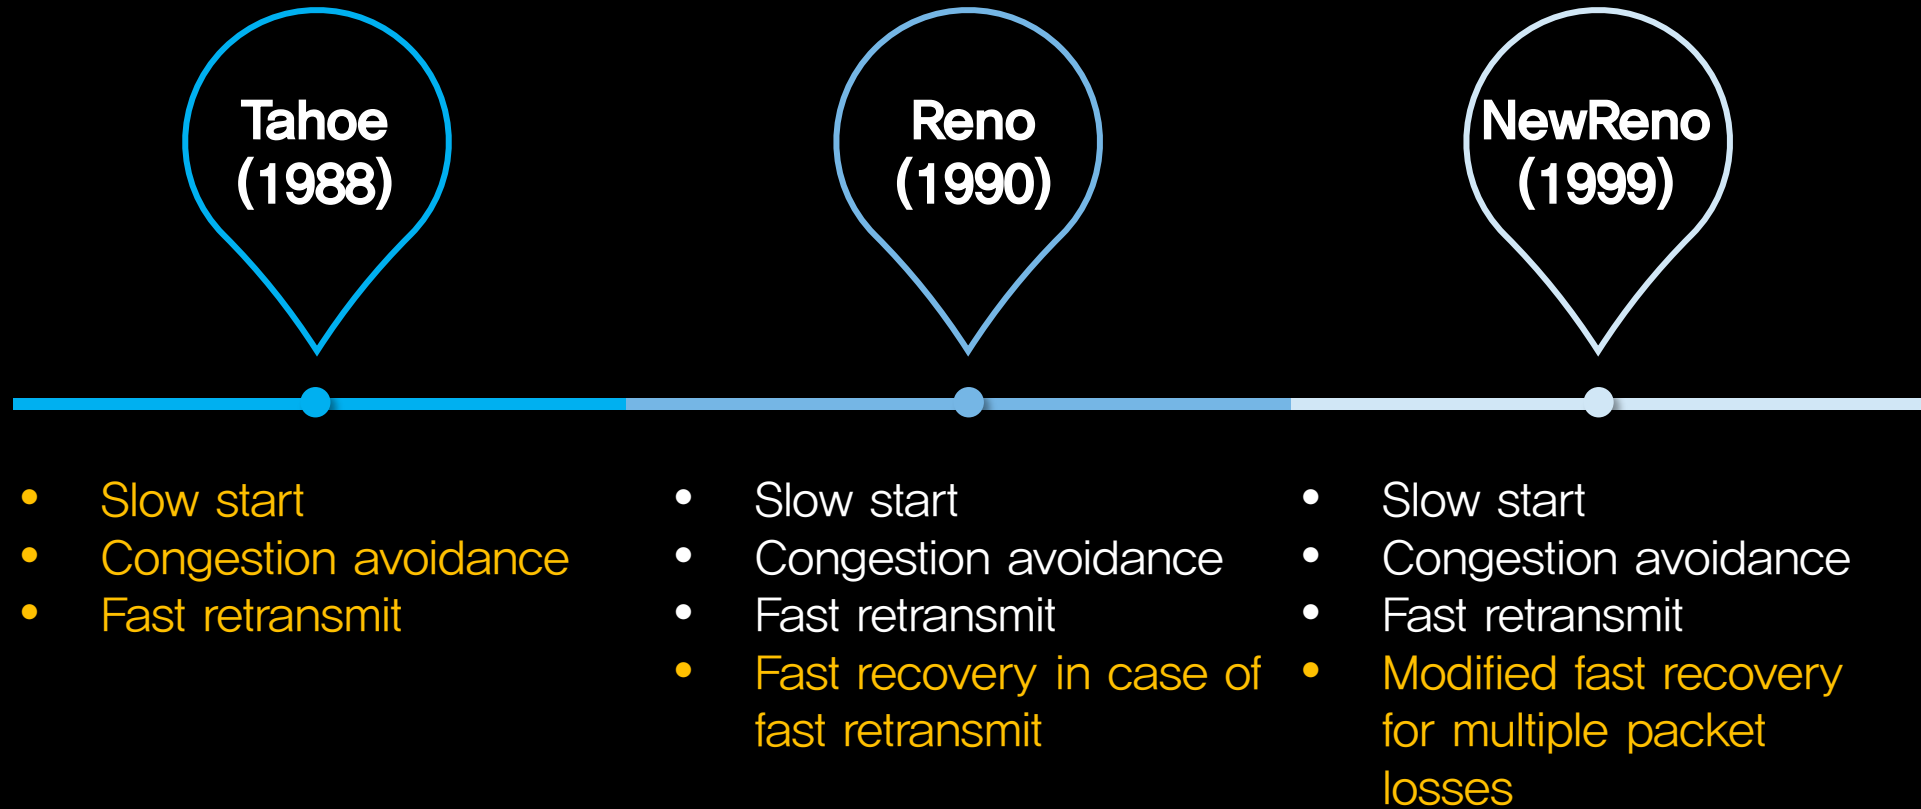
- Vulnerability of Reno's fast recovery
 - in case of multiple packet losses, **cwnd** may be reduced too much (exponential way)




출처 - A. Afanasyev et al., "Host-to-Host Congestion Control for TCP," IEEE Comm. Surveys & Tutorials, Vol. 12, No. 3, 3rd Quarter, 2010, pp. 304-342

- Refinement of Reno's fast recovery
 - resistant to multiple losses
 - restricts the exit from the recovery phase until all data packets from the initial congestion window are acknowledged
 - exits from the NewReno's fast recovery only when a new data ACK is received
 - keeps the sequence number of the last data packet sent before entering the fast recovery state

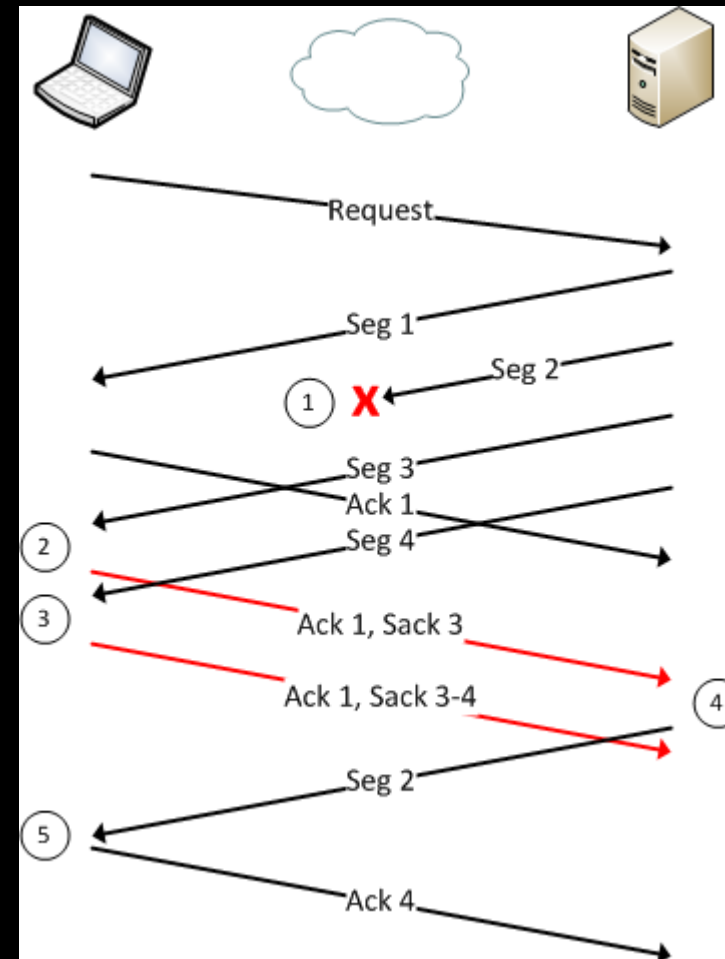




- Limited information of cumulative ACKs
 - ACK of only the last in-order packet
 - Reno's fast recovery assumes loss of only one data packet
 - a duration of the recovery is directly proportional to the number of packet losses
- **SACK (Selective ACK) option**
 - receiver provides information about several packet losses in a single ACK message by reporting blocks of successfully delivered data packets



Cumulative ACK



Selective ACK



08. TCP vs. UDP

- TCP fairness through AIMD (additive increase and multiplicative decrease)


- two competing sessions

- Multimedia apps often do not use TCP

- do not want rate throttled by congestion control

- Instead use UDP


- send audio/video at constant rate, tolerate packet loss





TCP

- **Slower but reliable transfers**
- **Typical applications:**
 - Email
 - Web browsing




unicast




UDP


- **Fast but non-guaranteed transfers (“best effort”)**
- **Typical applications:**
 - VoIP
 - Music streaming



unicast



multicast



broadcast

출처 -

https://www.google.co.kr/url?sa=i&rct=j&q=&esrc=s&source=images&cd=&cad=rja&uact=8&ved=2ahUKEwju_8aR0PXbAhVGf7wKHfEmAYUQJRx6BAgBEAU&url=https%3A%2F%2Fknowledgeofthings.com%2Ftcpip-vs-udp-internet-protocol-suite%2F&psig=AOvVaw2QofBlqkTFxG8_J4eyPGI&ust=1530250010414409/



Summary

01

Transport layer services

- program, process, thread
- logical communication between processes

02

Multiplexing and socket

- multiplexing and demultiplexing
- socket and socket program

03

User datagram protocol

- connectionless service
- fast but unreliable transfer

04

Reliable data transfer principles

- handling of bit error and packet loss
- ARQ: stop-and-wait, go-back-N, selective-repeat

05

Transmission control protocol

- connection-oriented service
- reliable data transfer through error control and flow control

06

Congestion control

- congestion: “Load is higher than network capacity!”
- additive increase and multiplicative decrease

07

TCP congestion control algorithm

- dynamics of congestion window
- TCP Tahoe, Reno, NewReno, SACK

08

TCP vs. UDP

- TCP: slower but reliable and fair
- UDP: fast but unguaranteed